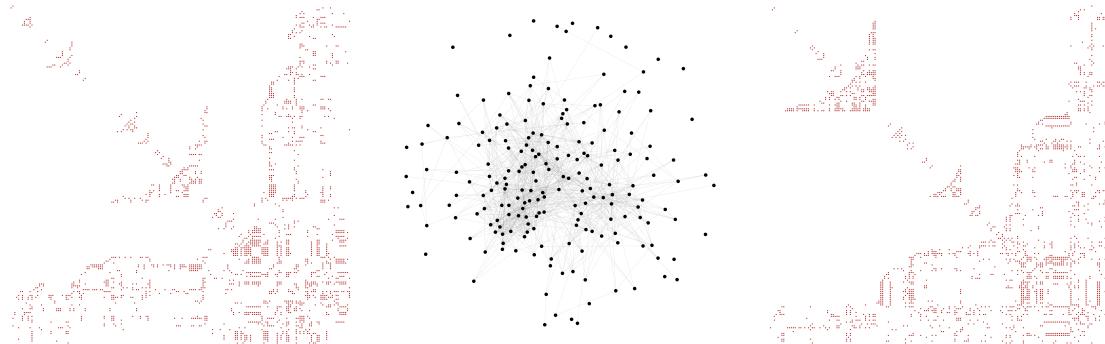


DEPARTMENT OF  
INFORMATION TECHNOLOGY AND ELECTRICAL ENGINEERING  
Fall Semester 2025

# Comprehensive Evaluation of Sparse Symmetric Matrix Permutation Methods

Master Thesis



Ritvik Ranjan  
[rranjan@ethz.ch](mailto:rranjan@ethz.ch)

October 1, 2025

Advisors: Vincent Maillou, [vmaillou@iis.ee.ethz.ch](mailto:vmaillou@iis.ee.ethz.ch)  
Dr. Alexandros Nikolaos Ziogas, [alziogas@iis.ee.ethz.ch](mailto:alziogas@iis.ee.ethz.ch)  
Professor: Prof. Dr. Mathieu Luisier, [leader@iis.ee.ethz.ch](mailto:leader@iis.ee.ethz.ch)

# Abstract

Sparse symmetric linear systems arise across computational science and engineering. Solving these systems efficiently via direct methods depends critically on matrix reordering to minimize fill-in during factorization. Finding optimal reorderings is NP-complete, necessitating good heuristic approaches. This thesis presents a comprehensive evaluation of sparse matrix reordering algorithms, comparing classical methods, graph partitioning approaches, and hypergraph-based techniques across diverse matrix structures. We assess multiple performance metrics including fill-in reduction, execution time, memory usage, and parallelization potential. Our benchmarking reveals that advanced methods achieve better fill-in reduction compared to classical algorithms, though at higher computational cost. We implemented several enhancements, including a Nested Dissection variant with improved coarsening and a GPU-accelerated RCM implementation, demonstrating significant performance improvements. The results establish some trade-offs between ordering quality and computational overhead, providing guidance for algorithm selection and future research directions in sparse matrix reordering.

# Acknowledgments

First and foremost, I would like to thank Prof. Dr. Mathieu Luisier for providing me with the opportunity to work in the Nano-TCAD group and on this fascinating topic.

I would like to thank Vincent Maillou for providing me the opportunity to work on this project and for his friendly and helpful supervision and the patience and time put into teaching me different aspects of it. I express my sincere thanks to Dr. Alexandros Nikolaos Ziogas for being supportive and guiding me through the process.



Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich

### Declaration of originality

The signed declaration of originality is a component of every written paper or thesis authored during the course of studies. **In consultation with the supervisor**, one of the following two options must be selected:

- I hereby declare that I authored the work in question independently, i.e. that no one helped me to author it. Suggestions from the supervisor regarding language and content are excepted. I used no generative artificial intelligence technologies<sup>1</sup>.
- I hereby declare that I authored the work in question independently. In doing so I only used the authorised aids, which included suggestions from the supervisor regarding language and content and generative artificial intelligence technologies. The use of the latter and the respective source declarations proceeded in consultation with the supervisor.

**Title of paper or thesis:**

Comprehensive Evaluation of Sparse Symmetric Matrix Permutation Methods

**Authored by:**

*If the work was compiled in a group, the names of all authors are required.*

**Last name(s):**

Ranjan

**First name(s):**

Ritvik

With my signature I confirm the following:

- I have adhered to the rules set out in the [Citation Guidelines](#).
- I have documented all methods, data and processes truthfully and fully.
- I have mentioned all persons who were significant facilitators of the work.

I am aware that the work may be screened electronically for originality.

**Place, date**

Zürich, 06.10.25

**Signature(s)**

*If the work was compiled in a group, the names of all authors are required. Through their signatures they vouch jointly for the entire content of the written work.*

<sup>1</sup> For further information please consult the ETH Zurich websites, e.g. <https://ethz.ch/en/the-eth-zurich/education/ai-in-education.html> and <https://library.ethz.ch/en/researching-and-publishing/scientific-writing-at-eth-zurich.html> (subject to change).

# Contents

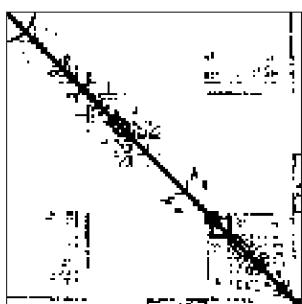
<b>1. Introduction</b>	<b>1</b>
<b>2. Background</b>	<b>3</b>
2.1. Elimination Trees and Symbolic Factorization . . . . .	3
2.1.1. Mathematical Foundations and Problem Formulation . . . . .	3
2.1.2. Structural Properties and Theoretical Characterizations . . . . .	5
2.1.3. Algorithmic Construction of Elimination Trees . . . . .	6
2.1.4. Symbolic Factorization . . . . .	6
2.2. Minimum fill-in and NP-completeness . . . . .	8
2.3. Heuristics classification . . . . .	9
2.3.1. Bandwidth Minimization . . . . .	9
2.3.2. Minimum Degree . . . . .	10
2.3.3. Nested Dissection . . . . .	10
2.3.4. Hypergraph partitioning and other methods . . . . .	11
<b>3. Implementation and Optimizations</b>	<b>12</b>
3.1. SuiteSparse implementation of RCM and Minimum Degree . . . . .	12
3.1.1. AMD Algorithm Overview . . . . .	12
3.1.2. COLAMD Algorithm Overview . . . . .	14
3.2. Nested Dissection using METIS and SCOTCH . . . . .	17
3.3. Parallel-Nested Dissection . . . . .	18
3.4. Parallelizing minimum degree . . . . .	19
3.5. New Coarsening approaches in Nested Dissection . . . . .	19
3.6. Hypergraph Based Ordering . . . . .	19
3.7. GPU Implementation of RCM . . . . .	19
<b>4. Results</b>	<b>22</b>
4.1. Evaluation setup . . . . .	22
4.1.1. Hardware Infrastructure . . . . .	22
4.2. Matrices dataset . . . . .	23

## *Contents*

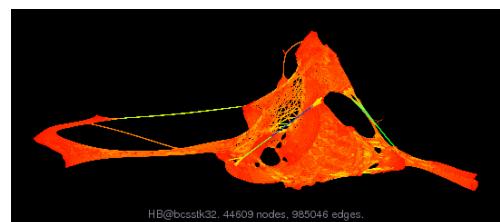
4.3. Methodology . . . . .	24
4.3.1. Performance Metrics and Measurement Protocol . . . . .	25
4.4. Evaluation . . . . .	26
4.5. Current limitations . . . . .	34
<b>5. Other Approaches</b>	<b>35</b>
5.1. Graph Reinforcement Learning for Reordering . . . . .	35
5.2. GPU Accelerated Nested Dissection . . . . .	36
<b>6. Conclusion and Future Work</b>	<b>37</b>
<b>A. Experimental Results</b>	<b>39</b>
<b>B. Task Description</b>	<b>45</b>
<b>List of Acronyms</b>	<b>46</b>
<b>List of Figures</b>	<b>47</b>
<b>List of Tables</b>	<b>49</b>
<b>Bibliography</b>	<b>50</b>

# Introduction

Sparse linear systems of the form  $Ax = b$ , where  $A \in \mathbb{R}^{n \times n}$  is sparse and symmetric positive definite, are ubiquitous in modern scientific computing. These systems arise naturally across virtually every domain of computational science and engineering, including quantum chemistry, computer graphics, computational fluid dynamics, power networks, machine learning, and optimization. Mathematically, a matrix  $A$  is considered sparse when the number of nonzero entries  $\text{nnz}(A)$  satisfies  $\text{nnz}(A) = O(n)$  rather than  $O(n^2)$  as in dense matrices.



(a) Sparsity pattern of the bcsstk32 matrix



(b) Graph representation

Figure 1.1.: The bcsstk32 matrix from automotive chassis analysis. This  $28,924 \times 28,924$  symmetric matrix represents an automobile chassis structure analyzed using finite element modeling.

This sparsity typically arises from the discretization of partial differential equations (PDEs) using finite difference, finite element, or finite volume methods. When continuous problems described by partial differential equations are discretized using methods such as finite differences, finite elements, or finite volumes, the resulting algebraic equations connect only neighboring (or some specific) grid points or elements. Concep-

## 1. Introduction

tually, sparsity corresponds to systems with few pairwise interactions. For instance, in a three-dimensional finite element mesh, each node typically connects to only a small neighborhood of adjacent nodes, regardless of the total problem size, resulting in matrices where the number of non-zero elements is roughly equal to the number of rows or columns rather than being proportional to  $n^2$ . This is shown in an example in Fig. 1.1, which illustrates the sparsity pattern of a large matrix arising from automotive chassis analysis, or Fig. 1.2 which shows a quantum transport device.

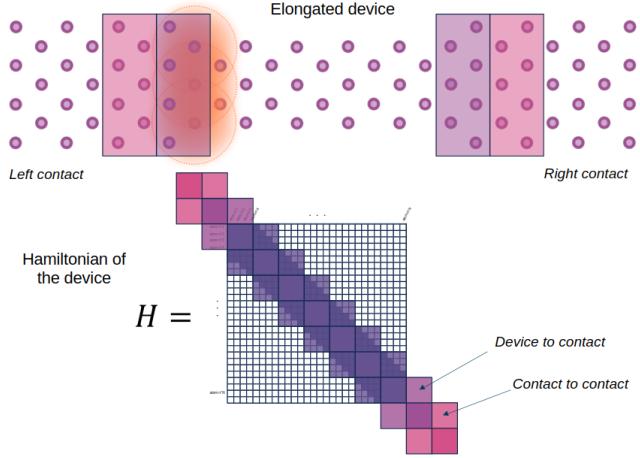


Figure 1.2.: Sparse matrix structure from a quantum transport device. The block-tridiagonal pattern reflects local physical interactions between neighboring spatial layers.

Matrix permutations serve as a preprocessing step that can dramatically improve both computational efficiency and memory requirements for sparse linear system solution. The basic principle involves reordering the rows and columns of the matrix  $A$  to obtain a permuted system  $PAP^T \hat{x} = Pb$ , where  $P$  is a permutation matrix, such that the reordered matrix exhibits superior properties for factorization. The primary motivation for permutation is fill-in reduction during matrix factorization. Fill-in occurs when zeros in the original matrix become non-zero during factorization, effectively destroying the sparse structure.

In this thesis, we focus exclusively on the permutation problem for sparse symmetric matrices. We explore existing methods for sparse symmetric matrix reordering, investigate potential optimizations and parallelization strategies, and evaluate the performance of different reordering methods across a variety of matrices, assessing their effectiveness in terms of fill-in reduction, memory usage and parallelism. In chapter 2, we provide the necessary theoretical background on sparse matrix factorization and graph representations. Chapter 3 details the implementation of various reordering algorithms, while chapter 4 presents an overview of our experimental evaluation. For full performance results, please refer to the appendix A.

# Chapter 2

## Background

Before we look into the algorithms and theory behind reordering techniques, we first take a dive into the factorization process that produces further non-zeroes, which in this and the entire thesis's context, are called fill-in. Fill-in depends only on the structure of the matrix, i.e., the positions where the initial non-zero entries are placed. It may be trivial but important to note, that if some numerical arithmetic results in a zero in the factored matrix, it is usually still considered as a fill-in.

Suppose the set of set of  $n$  equations that are to be solved are

$$Ax = b \tag{2.1}$$

where  $A$  is a sparse matrix,  $x$  is the vector of unknowns. The sparsity, or the number of non-zero entries in  $A$  determines the fill-in of its Cholesky factor when it is employed to solve the aforementioned set of equations. Suppose the Cholesky factorization of  $A$  is given by  $LL^T$ , where  $L$  is a lower triangular matrix (with a positive diagonal) and  $L^T$  is the transpose of  $L$ . The efficiency of solving this set of equations is dependent on the number of non-zero entries in  $L$ . It has been shown (George and Liu) that we can factor and solve the set of equations in space proportional to  $\sum_j d_j$  and time complexity  $\sum_j d_j j^2$ , where  $d_j$  is the number of non-zeroes in the  $j$ th column of  $L$ .

### 2.1. Elimination Trees and Symbolic Factorization

#### 2.1.1. Mathematical Foundations and Problem Formulation

##### Sparse Matrix Factorization Context

Consider a sparse symmetric positive definite matrix  $A \in \mathbb{R}^{n \times n}$ . The Cholesky factorization decomposes  $A$  into the form  $A = LL^T$ , where  $L$  is a lower triangular matrix. While  $A$  may contain relatively few nonzero entries, the factor  $L$  typically exhibits significantly more nonzeros due to a phenomenon known as fill-in. This fill-in occurs because the

## 2. Background

elimination process creates new nonzero entries at positions that were originally zero in  $A$ .

The sparsity pattern of  $A$  is defined as the set  $\text{pattern}(A) = \{(i,j) : A_{ij} \neq 0\}$ . Similarly,  $\text{pattern}(L)$  denotes the sparsity pattern of the Cholesky factor. The fundamental observation is that  $\text{pattern}(L)$  can be completely determined from  $\text{pattern}(A)$  using purely structural analysis, without requiring any numerical computation. This structural predictability forms the theoretical foundation for elimination trees and symbolic factorization.

### Graph-Theoretic Representation of Sparse Matrices

Every sparse symmetric matrix  $A$  can be naturally represented as an undirected graph  $G(A) = (V, E)$ , where the vertex set  $V = \{1, 2, \dots, n\}$  corresponds to the rows and columns of  $A$ , and the edge set  $E$  contains an edge  $(i, j)$  if and only if  $A_{ij} \neq 0$  for  $i \neq j$ . This graph representation transforms matrix operations into graph-theoretic problems, enabling the application of powerful combinatorial methods to numerical linear algebra.

The elimination process on matrix  $A$  corresponds to a sequence of vertex eliminations on graph  $G(A)$ . When vertex  $k$  is eliminated, all pairs of neighbors of  $k$  become connected by edges if they were not already connected. This edge addition process continues throughout the elimination sequence, producing what is known as the filled graph  $G^+(A)$ . The filled graph contains all edges that exist at any point during the elimination process, and its structure completely determines the sparsity pattern of the Cholesky factor  $L$ .

### Formal Definition of Elimination Trees

The elimination tree provides a compact representation of the structural relationships inherent in sparse matrix factorization. This tree structure captures the essential dependencies between different stages of the elimination process.

**Definition 2.1** (Elimination Tree). Let  $A$  be a symmetric positive definite matrix with Cholesky factorization  $A = LL^T$ . The elimination tree  $T(A) = (V, E_T)$  is a directed tree defined by the parent function  $\text{parent} : V \rightarrow V \cup \{\emptyset\}$ , where:

$$\text{parent}(j) = \min\{i > j : L_{ij} \neq 0\} \quad (2.2)$$

If no such index  $i$  exists, then  $\text{parent}(j) = \emptyset$  and  $j$  is a root of the tree. The directed edges of the tree are given by  $E_T = \{(j, \text{parent}(j)) : \text{parent}(j) \neq \emptyset\}$ .

This definition establishes a precise correspondence between the numerical structure of the Cholesky factor and the combinatorial structure of a directed tree. Each node  $j$  in the tree corresponds to column  $j$  of the matrix, and the parent relationship encodes which column will first modify column  $j$  during the factorization process.

**Definition 2.2** (Ancestor Relationship). In the elimination tree  $T(A)$ , node  $i$  is an ancestor of node  $j$  (denoted  $i \succ j$ ) if there exists a directed path from  $j$  to  $i$  in the tree. Node  $i$  is a proper ancestor if  $i \succ j$  and  $i \neq j$ .

## 2. Background

The ancestor relationship in the elimination tree has a profound connection to the structure of the Cholesky factor, as formalized in the following fundamental theorem.

**Theorem 2.1** (Fundamental Characterization Theorem). *Let  $T(A)$  be the elimination tree of symmetric positive definite matrix  $A$  with Cholesky factor  $L$ . Then for any indices  $i > j$ :  $L_{ij} \neq 0$  if and only if  $i \succ j$  in  $T(A)$ .*

This theorem establishes that the elimination tree completely characterizes the sparsity pattern of the Cholesky factor. The nonzero entries in column  $j$  of  $L$  correspond exactly to the ancestors of node  $j$  in the elimination tree. This remarkable result shows that purely combinatorial information (tree ancestry) determines numerical structural properties (matrix sparsity patterns).

### 2.1.2. Structural Properties and Theoretical Characterizations

#### Fill-in Characterization Through Path Analysis

Understanding when fill-in occurs during elimination requires analyzing connectivity patterns in the original matrix graph. The concept of fill paths provides a precise characterization of this phenomenon.

**Definition 2.3** (Fill Path). Let  $G(A) = (V, E)$  be the graph of matrix  $A$ . A fill path from vertex  $i$  to vertex  $j$  (where  $i < j$ ) is a path  $P = (i = v_0, v_1, v_2, \dots, v_k = j)$  in  $G(A)$  such that all intermediate vertices satisfy  $v_l < \min(i, j)$  for  $1 \leq l \leq k - 1$ .

The fill path condition ensures that all intermediate vertices on the path are eliminated before either endpoint, which is precisely the condition under which the elimination process will create a direct connection between  $i$  and  $j$ .

**Theorem 2.2** (Fill Path Characterization). *An edge  $(i, j)$  with  $i < j$  belongs to the filled graph  $G^+(A)$  if and only if there exists a fill path from  $i$  to  $j$  in the original graph  $G(A)$ .*

This theorem provides the theoretical foundation for predicting fill-in patterns. It shows that structural properties of the original graph completely determine which new edges will be created during elimination, independent of the specific numerical values in the matrix.

#### Postorder Properties and Structural Invariants

The elimination tree possesses several important structural properties that make it amenable to efficient algorithmic manipulation.

**Theorem 2.3** (Postorder Property). *In any elimination tree  $T(A)$ , if  $\text{parent}(j) = i$ , then all nodes in the subtree rooted at  $j$  have indices less than  $i$ . Furthermore, the elimination tree is uniquely determined by the sparsity pattern of  $A$ , independent of the specific elimination ordering used (provided the ordering respects the natural ordering of indices).*

## 2. Background

This property implies that elimination trees can be processed using postorder traversal, where children are visited before their parents. This traversal order naturally reflects the dependencies in the elimination process—a column can only be processed after all columns in its subtree have been completed.

### 2.1.3. Algorithmic Construction of Elimination Trees

The most straightforward approach to constructing elimination trees directly implements the definition by examining the sparsity structure of the matrix.

---

**Algorithm 1:** Direct Construction Method

---

**Input :**Symmetric positive definite matrix  $A$  with sparsity pattern  
**Output:**Parent array representing elimination tree  $T(A)$

```
1 Initialize parent[1...n] ← 0;                                // 0 represents null parent
2 for  $j = 1$  to  $n$  do
3   parent[j] ← 0;
4   foreach  $i \in \{k : A[j,k] \neq 0 \text{ and } k > j\}$  do
5     if parent[j] = 0 or  $i < \text{parent}[j]$  then
6       parent[j] ← i;
7     end
8   end
9 end
10 return parent array
```

---

This algorithm requires  $O(|A|)$  time, where  $|A|$  denotes the number of nonzero entries in matrix  $A$ . The space complexity is  $O(n)$  for storing the parent array. While simple to understand and implement, this direct approach does not exploit the structural properties of elimination trees for potential efficiency improvements.

### 2.1.4. Symbolic Factorization

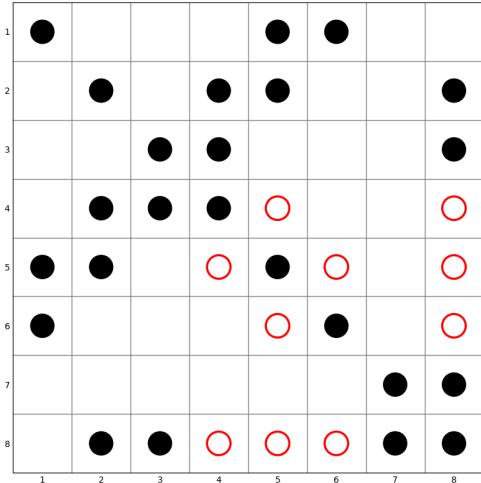
Symbolic factorization determines the sparsity pattern of the Cholesky factor  $L$  from the sparsity pattern of matrix  $A$  using the elimination tree. It does not involve any numerical computations and relies solely on the structure of the original matrix.

The key insight is that the sparsity structure of each column in  $L$  can be computed by analyzing how structural information propagates through the elimination tree. Each column inherits structure from two sources: the original matrix  $A$  and the previously computed columns corresponding to its children in the elimination tree.

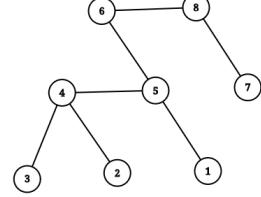
**Definition 2.4** (Column Structure Inheritance). For column  $j$  of the Cholesky factor  $L$ , the nonzero pattern is determined by:

- Direct inheritance: All nonzero entries from column  $j$  of the original matrix  $A$  that lie on or below the diagonal

## 2. Background



(a) Sparse matrix structure



(b) Corresponding elimination tree

Figure 2.1.: Illustration of elimination tree construction for the given sparse symmetric matrix. The sparsity is denoted by filled black circles and the fill-in induced is denoted by hollow red circles.

- Indirect inheritance: All entries from the union of nonzero patterns of columns corresponding to children of  $j$  in the elimination tree, restricted to rows numbered greater than  $j$

This inheritance pattern reflects the computational dependencies during numerical factorization—column  $j$  can only be processed after all its children in the elimination tree have been completed.

The postorder traversal ensures that each column is processed only after all its descendants in the elimination tree have been completed, respecting the computational dependencies inherent in the factorization process.

**Theorem 2.4** (Correctness of Symbolic Factorization). *Algorithm 2 correctly computes the sparsity pattern of the Cholesky factor  $L$  in  $O(|L|)$  time, where  $|L|$  denotes the number of nonzero entries in  $L$ .*

## 2. Background

---

**Algorithm 2:** Elimination Tree Symbolic Factorization

---

**Input** :Sparsity pattern of matrix  $A$ , elimination tree  $T(A)$   
**Output**:Sparsity pattern of Cholesky factor  $L$

```

1 Initialize column_structure[1...n] as empty sets;
2 for  $j = 1$  to  $n$ ;                                // in postorder traversal of  $T(A)$ 
3   do
4     // Direct inheritance from original matrix
5     column_structure[j]  $\leftarrow \{i \geq j : A[i,j] \neq 0\}$ ;
6     // Indirect inheritance from children in elimination tree
7     foreach child  $c$  of  $j$  in  $T(A)$  do
8       | column_structure[j]  $\leftarrow$  column_structure[j]  $\cup \{i \in \text{column\_structure}[c] :$ 
9       |    $i > j\}$ ;
10      end
11      // Record structure for column  $j$  of  $L$ 
12      pattern( $L[:,j]$ )  $\leftarrow$  column_structure[j];
13   end
14 return pattern( $L$ )

```

---

### 2.2. Minimum fill-in and NP-completeness

When performing Cholesky factorization on a sparse symmetric matrix, we want to minimize fill-in - the new nonzero entries that appear during elimination. This translates to a graph theory problem: given a graph representing the matrix's sparsity pattern, find an elimination order that creates the fewest new edges. Yannakakis [1] proved that this optimization problem is NP-complete, meaning no polynomial-time algorithm can solve it optimally in general (unless P = NP).

The key insight connecting linear algebra to graph theory is that chordal graphs are essential in proving this. A graph is chordal if every cycle of length  $\geq 4$  has a chord (an edge connecting non-consecutive vertices). The fundamental theorem states that a graph has perfect elimination (zero fill-in) if and only if it is chordal, and any elimination process on any graph produces a chordal result. Therefore, the minimum fill-in problem becomes: What is the smallest number of edges we must add to make our graph chordal?

The proof uses a three-step reduction to prove NP-completeness. First, we consider bipartite graphs and their special case, chain graphs. A bipartite graph has vertices split into two independent sets  $P$  and  $Q$ , and it is a chain graph when neighborhoods form a nested sequence:  $\Gamma(v_1) \subseteq \Gamma(v_2) \subseteq \dots \subseteq \Gamma(v_k)$ . The key property is that a bipartite graph is a chain graph if and only if it contains no pair of independent edges.

Second, for any bipartite graph  $G = (P, Q, E)$ , we construct  $C(G)$  by making  $P$  into a complete clique (all vertices in  $P$  connected), making  $Q$  into a complete clique (all vertices in  $Q$  connected), and keeping original edges between  $P$  and  $Q$ . The crucial lemma states that  $C(G)$  is chordal if and only if  $G$  is a chain graph. This works because

## 2. Background

if  $G$  has independent edges, they create a 4-cycle in  $C(G)$  with no chord. Conversely, if  $G$  is a chain graph,  $C(G)$  has a perfect elimination order.

Third, the reduction uses the Optimal Linear Arrangement Problem, which asks: given a graph, arrange vertices on a line to minimize the sum of distances between adjacent vertices. This problem is known to be NP-complete. Yannakakis constructs a transformation where, given graph  $G$  with  $n$  vertices and  $m$  edges, he creates bipartite graph  $G'$  where part  $P$  has one vertex for each vertex in  $G$ , and part  $Q$  has an elaborate gadget structure with  $2m + n^2 - d(v)$  vertices, with connections that encode the linear arrangement constraints. The mathematical relationship is:

$$\text{Minimum fill-in of } G' = \text{Optimal arrangement cost of } G + \frac{n^2(n-1)}{2} - 2m \quad (2.3)$$

This result has practical implications that no perfect algorithm exists, as any algorithm guaranteeing optimal fill-in will require exponential time in the worst case. This explains why practical sparse matrix software uses heuristics like minimum degree ordering, nested dissection and such.

### 2.3. Heuristics classification

There are several heuristics that have been proposed to reduce the fill-in developed over the years. These heuristics can be broadly classified in the following categories: Bandwidth minimization, minimum degree (and its variants), nested dissection, banded structure methods using hypergraphs, and recently, machine learning approaches.

#### 2.3.1. Bandwidth Minimization

Bandwidth minimization refers to the problem of permuting the rows and columns of a matrix such that the non-zero entries are as close to the diagonal as possible. Gaussian elimination can be performed in  $O(nb^2)$  time on matrices of dimension  $n$  and bandwidth  $b$ , which is faster than the forward  $O(n^3)$  algorithm when  $b$  is smaller than  $n$  (Lim A. et al., 2006a). Additionally, this problem is NP-complete, but several heuristics have been developed to approximate the optimal solution.

The Cuthill-McKee algorithm employs breadth-first search to reduce matrix bandwidth by generating level structures. However, it has computational limitations and may not achieve optimal bandwidth reduction. The Reverse Cuthill-McKee algorithm addresses these issues by reversing the ordering, while the GPS algorithm by Gibbs et al. provides an alternative level structure approach.

We take a look at the reverse Cuthill-McKee algorithm, which is still widely used in practice.

[write the pseudocode here]

## 2. Background

### 2.3.2. Minimum Degree

The minimum degree algorithm determines the order in which to eliminate variables (or pivot) during Gaussian elimination to minimize fill-in (creation of new non-zero entries) in sparse matrices. At each step, it chooses the node with the minimum degree (fewest connections) for elimination.

The minimum degree algorithm operates on the adjacency graph of the sparse matrix. It begins by initializing the graph, where each vertex represents a variable. At each iteration, the algorithm selects the vertex with the smallest degree (i.e., the fewest neighbors), which corresponds to the variable whose elimination is expected to introduce the least fill-in. Upon elimination, the chosen vertex is removed from the graph, and all its neighbors are connected to each other, forming a clique to preserve the matrix structure. The degrees of the affected vertices are then updated to reflect the new connections. This process is repeated until all vertices have been eliminated. The resulting elimination order directly determines the pivot sequence used during matrix factorization.

(placeholder for general algorithm)

Variants of the minimum degree algorithm include:

#### 1. Multiple Minimum Degree (MMD)

When multiple nodes have the same minimum degree, uses additional tie-breaking rules. Often selects based on secondary criteria like external degree or node numbering.

#### 2. Approximate Minimum Degree (AMD)

Uses approximations to avoid expensive exact degree calculations. Employs concepts like "supervariables" and "element absorption". Much faster than exact minimum degree while maintaining similar fill-in quality.

#### 3. Constrained Minimum Degree

Incorporates additional constraints (like maintaining bandwidth). Balances minimum degree objective with other structural requirements.

### 2.3.3. Nested Dissection

Nested dissection is a divide and conquer heuristic for the solution of sparse symmetric systems of linear equations based on graph partitioning. Nested dissection can be viewed as a recursive divide-and-conquer algorithm on an undirected graph; it uses separators in the graph, which are small sets of vertices whose removal divides the graph approximately in half [2].

Nested dissection consists of the following steps: Form an undirected graph in which the vertices represent rows and columns of the system of linear equations, and an edge represents a nonzero entry in the sparse matrix representing the system. Recursively

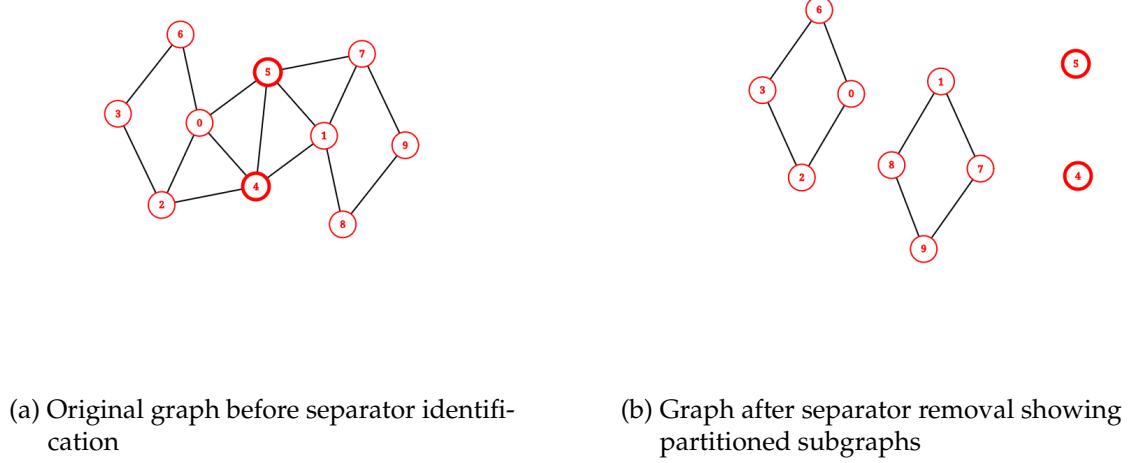


Figure 2.2.: Separator identification process. The separator vertices divide the graph into approximately equal subgraphs

partition the graph into subgraphs using separators, small subsets of vertices the removal of which allows the graph to be partitioned into subgraphs with at most a constant fraction of the number of vertices. Perform Cholesky decomposition (a variant of Gaussian elimination for symmetric matrices), ordering the elimination of the variables by the recursive structure of the partition: each of the two subgraphs formed by removing the separator is eliminated first, and then the separator vertices are eliminated [3].

All the sequential algorithms for determining the elimination ordering of a graph  $G$  can be described by the following general algorithm: Generate the tree of separators for  $G$  and perform a tree traversal on the separator tree to order the vertices, where this traversal must visit a node before any of its parents.

### 2.3.4. Hypergraph partitioning and other methods

In recent years, hypergraph partitioning methods [4] have been proposed to reduce fill-in during sparse matrix factorization. Hypergraphs generalize graphs by allowing edges (hyperedges) to connect more than two vertices, making them well-suited for modeling complex relationships in sparse matrices. In this paper we explore another Hypergraph based reordering method,

Other methods include machine learning approaches that learn optimal ordering strategies using reinforcement learning. One such approach is Alpha Elimination [5] which uses Convolutional Neural Networks (CNNs) and reinforcement learning to predict the next node to eliminate based on the current state of the graph. We have explored a more flexible architecture but fails with scalability, which we discuss in Chapter 5.

# Chapter 3

## Implementation and Optimizations

Discuss the state-of-the art and other related work. Depending on how much related work exists and how central the comparison to it is in your thesis (discuss this with your advisors), the introduction may contain sufficient related work and this chapter can be omitted.

### 3.1. SuiteSparse implementation of RCM and Minimum Degree

The Reverse Cuthill-McKee (RCM) algorithm in SuiteSparse provides bandwidth reduction for symmetric sparse matrices through a breadth-first search strategy combined with degree-based ordering heuristics. The implementation processes disconnected components separately and employs pseudo-peripheral vertex selection to minimize profile and bandwidth.

The SuiteSparse RCM implementation incorporates several refinements over the basic algorithm. The pseudo-peripheral vertex selection uses multiple BFS traversals to identify vertices that are approximately diametrically opposite, which typically results in better bandwidth reduction than arbitrary starting points. The degree-based sorting of neighbors during BFS traversal helps create a more systematic ordering that tends to group low-degree vertices together, further improving the resulting bandwidth.

The algorithm's effectiveness stems from its ability to produce orderings where vertices with similar connectivity patterns are placed close together in the permutation. This locality property translates directly into reduced bandwidth and improved cache performance during matrix operations, making RCM particularly valuable for iterative solvers and direct factorization methods that benefit from band structure preservation.

#### 3.1.1. AMD Algorithm Overview

The Approximate Minimum Degree (AMD) algorithm implemented in SuiteSparse follows a refined elimination-based approach that balances computational efficiency with fill-in minimization. The algorithm operates on a quotient graph representation and

### 3. Implementation and Optimizations

---

**Algorithm 3:** SuiteSparse RCM Algorithm

---

```

input :Symmetric matrix  $A$  ( $n \times n$ )
output:Permutation  $P$  for bandwidth reduction

1 Initialize;;
2 Compute degree[ $i$ ] = |adj( $i$ )| for all vertices  $i$ ;
3 Find connected components  $R_1, R_2, \dots, R_k$  using DFS;
4 Initialize visited[ $i$ ] = false for all  $i$ ;
5 Set perm_index = 0;
6 Process each component;;
7 foreach connected component  $R_j$  do
8   Find pseudo-peripheral start vertex;;
9   start = FindPseudoPeripheral( $R_j$ );
   // Select vertex with minimum degree among maximum-distance
   // vertices
10  Cuthill-McKee BFS traversal;;
11  Initialize queue  $Q = \{\text{start}\}$ ;
12  Set visited[start] = true;
13  Initialize CM_order = [start];
14  while  $Q \neq \emptyset$  do
15     $u = \text{dequeue}(Q)$ ;
16    neighbors = sort(unvisited_adj( $u$ ), by_degreeAscending);
17    foreach  $v \in \text{neighbors}$  do
18      if  $\neg\text{visited}[v]$  then
19        Set visited[v] = true;
20        enqueue( $Q, v$ );
21        Append  $v$  to CM_order;
22      end
23    end
24  end
25 Apply reverse ordering (RCM);;
26 for  $i = 0$  to  $|\text{CM\_order}| - 1$  do
27    $| P[\text{perm\_index} + i] = \text{CM\_order}[|\text{CM\_order}| - 1 - i]$ ;
28 end
29  $\text{perm\_index} \leftarrow \text{perm\_index} + |\text{CM\_order}|$ ;
30 end
31 return  $P$ ;

```

---

### 3. Implementation and Optimizations

incorporates several key optimizations including aggressive absorption, approximate degree updates, and dense row detection.

The key innovation in SuiteSparse's AMD implementation lies in its aggressive absorption strategy and approximate degree computation. The aggressive absorption phase identifies elements that can be completely absorbed into the current pivot, reducing the size of the quotient graph and improving cache locality. The approximate degree updates provide a computationally efficient method to maintain ordering decisions without exact degree computation, which becomes prohibitively expensive as elimination progresses.

The dense row detection mechanism addresses a common pathology in minimum degree algorithms where elimination of high-degree vertices can lead to excessive fill-in. When the algorithm detects that a newly formed element exceeds a threshold based on the matrix size, it defers elimination of the associated variables, effectively implementing a hybrid strategy that combines minimum degree with nested dissection principles.

#### 3.1.2. COLAMD Algorithm Overview

The Column Approximate Minimum Degree (COLAMD) algorithm in SuiteSparse extends the minimum degree concept to rectangular matrices by operating on a bipartite graph representation. Unlike AMD which works on the symmetric structure  $A^T A$ , COLAMD directly processes the rectangular matrix  $A$  to produce a column ordering that minimizes fill-in during factorization.

The COLAMD algorithm addresses the unique challenges of rectangular matrix ordering by maintaining both row and column degree information throughout the elimination process. The bipartite graph formulation allows the algorithm to track how column eliminations affect the sparsity structure without explicitly forming the potentially much denser  $A^T A$  matrix. The dense row detection mechanism prevents pathological behavior when processing matrices with highly dense rows, which could otherwise lead to excessive fill-in during factorization.

The degree update strategy in COLAMD carefully accounts for the fill-in patterns specific to rectangular matrices, where eliminating a column affects all other columns that share nonzero entries in the same rows. The algorithm's ability to handle supernodes (groups of columns with identical sparsity patterns) further enhances its effectiveness for structured matrices commonly arising in finite element applications.

### 3. Implementation and Optimizations

---

**Algorithm 4:** SuiteSparse AMD Algorithm

---

**input** :Symmetric matrix  $A$  ( $n \times n$ ), Control parameters  
**output**:Permutation  $P$ , Info statistics

```

1 Initialize:;
2 Convert  $A$  to  $A + A^T$  pattern if unsymmetric;
3 Build quotient graph  $G = (V, E)$  from  $A$ ;
4 Initialize degree lists Head[ $d$ ] for  $d = 0$  to  $n$ ;
5 Set degree[ $v$ ] =  $|\text{adj}(v)|$  for all vertices  $v$ ;
6 Place each vertex in appropriate degree list;

7 Main elimination loop:;
8 for  $k = 1$  to  $n$  do
9   Select pivot:;
10  Find minimum degree  $d$  with non-empty Head[ $d$ ];
11  Select pivot  $p$  from Head[ $d$ ];
12  Remove  $p$  from degree list;
13  Set  $P[k] = p$  (add to elimination ordering);
14  Element absorption:;
15  foreach element  $e$  adjacent to  $p$  do
16    if  $|L_e \cap L_p| = |L_e|$  then
17      | Absorb  $e$  into  $p$  (aggressive absorption);
18    end
19  end
20  Form new element  $e_p$ :;
21   $L_{e_p} = \text{adj}(p) \setminus \{\text{absorbed elements}\}$ ;
22  Mark  $e_p$  as new element;
23  Update degrees (approximate):;
24  foreach uneliminated vertex  $v \in L_{e_p}$  do
25    external_degree[ $v$ ] =  $|\text{adj}(v) \cap \text{uneliminated}|$ ;
26    bound =  $|L_{e_p}| - |L_{e_p} \cap \text{adj}(v)|$ ;
27    degree[ $v$ ]  $\approx$  external_degree[ $v$ ] + bound;
28    Move  $v$  to new degree list;
29  end
30  Dense row detection:;
31  if  $|L_{e_p}| > \max(\alpha\sqrt{n}, 16)$  then
32    | Mark  $e_p$  as dense element;
33    | Move dense variables to end of ordering;
34  end
35 end
36 Post-processing:;
37 Apply elimination tree post-ordering;
38 Compute final permutation statistics;
39 return  $P$  and Info;

```

---

### 3. Implementation and Optimizations

---

**Algorithm 5:** SuiteSparse COLAMD Algorithm

---

**input :**Matrix  $A$  ( $m \times n$ ) in CSC format, Control parameters  
**output:**Column permutation  $P$ , Info statistics

1 *Initialize*;  
2 Treat  $A$  as bipartite graph  $G = (R \cup C, E)$ ;  
3 Set  $\text{col\_degree}[j] = \text{nonzeros in column } j$ , place in degree list;  
4 Detect dense rows: if  $\text{row\_degree}[i] > \text{threshold}$ , mark dense;  
5 *Main elimination loop*;  
6 **for**  $k = 1$  **to**  $n$  **do**  
7     Select minimum degree column  $c$  (tie-break by density);  
8     Set  $P[k] = c$ ;  
9     *Update degrees*;  
10    **foreach** row  $i$  connected to  $c$  **do**  
11      **foreach** uneliminated column  $j$  in row  $i$  **do**  
12        fill\_count =  $|R(c) \cap R(j)|$ ;  
13        new\_degree =  $\text{col\_degree}[j] + |R(c)| - \text{fill\_count} - 1$ ;  
14        Apply dense penalty if row  $i$  is dense;  
15        Move  $j$  to Head[new\_degree];  
16      **end**  
17    **end**  
18    Mark eliminated rows and handle supernode formation;  
19 **end**  
20 Place remaining dense columns at end of ordering;  
21 **return**  $P$  and Info;

---

### 3. Implementation and Optimizations

#### 3.2. Nested Dissection using METIS and SCOTCH

Nested dissection represents a divide-and-conquer approach to matrix ordering that recursively partitions the graph using small vertex separators, ordering the separated components before the separator vertices. METIS and SCOTCH implement sophisticated multilevel nested dissection algorithms that combine graph coarsening, separator finding, and refinement techniques to produce high-quality orderings for large sparse matrices.

The multilevel nested dissection approach in METIS provides superior ordering quality compared to single-level methods by operating at multiple scales. The coarsening phase creates a hierarchy of increasingly smaller graphs while preserving essential structural properties through heavy edge matching. This matching strategy prioritizes edges with large weights, which in the context of matrix ordering typically correspond to strong structural connections that should be preserved during coarsening.

The separator computation on the coarsest level benefits from reduced problem size while maintaining global structural awareness. The refinement phase during projection ensures that separators remain high-quality as they are mapped back to finer graph levels. The recursive application of this process creates a natural hierarchy where large components are isolated first, followed by progressively smaller substructures, resulting in elimination orderings with excellent fill-in characteristics for sparse direct solvers.

---

**Algorithm 6:** Multilevel Graph Coarsening

---

```
input :Graph  $G = (V, E)$ 
output: Hierarchy of progressively coarser graphs
1 Hierarchy = [G] // Start with original graph
2 CurrentGraph = G;
3 while  $|V(CurrentGraph)| > COARSENING_THRESHOLD$  do
4   Find heavy edge matching to preserve graph structure;
5   Matching = HeavyEdgeMatching(CurrentGraph);
6   Contract matched edges to create coarser graph;
7   CoarserGraph = ContractEdges(CurrentGraph, Matching);
8   Add to hierarchy;
9   Hierarchy.append(CoarserGraph);
10  CurrentGraph = CoarserGraph;
11 end
12 return Hierarchy;
```

---

### 3. Implementation and Optimizations

---

**Algorithm 7:** Heavy Edge Matching Algorithm

---

```
input :Graph G with edge weights
output:Maximal matching favoring heavy edges

1 Matching = {};
2 Matched = {} // Track matched vertices
3 Sort edges by weight, centralness and degree for better matching quality;
4 SortedEdges = SortByWeight(E(G),descending = true);
5 foreach edge (u,v) in SortedEdges do
6   | if u ∉ Matched and v ∉ Matched then
7   |   | Matching.add((u,v));
8   |   | Matched.add(u);
9   |   | Matched.add(v);
10  | end
11 end
12 return Matching;
```

---

### 3.3. Parallel-Nested Dissection

### *3. Implementation and Optimizations*

#### **3.4. Parallelizing minimum degree**

There haven't been many attempts to parallelize the minimum degree algorithm due to its inherently sequential nature. The only known approximate parallel implementation of the minimum degree algorithm is the ParAMD algorithm proposed by Chang et al. in [6].

The sequential AMD algorithm has inherent bottlenecks that make parallelization difficult. Each elimination step requires selecting a pivot with minimum approximate degree, after which the degrees of neighboring variables must be updated. These steps are inherently sequential since you cannot select the next pivot until all updates from the previous elimination are complete.

Instead of eliminating one pivot at a time, the algorithm selects multiple pivots simultaneously using "distance-2 independent sets" - pivots that are at least 2 steps apart in the graph. This ensures no overlap in pivot neighborhoods, eliminating contention between parallel threads.

The algorithm allows selection of pivots whose degrees are within a multiplicative factor (`mult`) of the minimum degree. This relaxation increases the pool of available pivots for parallel processing while balancing against ordering quality - too much relaxation degrades the solution.

The ParAMD algorithm employs concurrent connection updates by pre-allocating 1.5x the original graph storage to avoid dynamic memory allocation. Each thread claims space atomically after collecting all updates, eliminating garbage collection synchronization bottlenecks.

For degree management, each thread maintains its own degree lists instead of sharing a global structure. The algorithm uses an affinity array to track which thread has the most current information for each variable, with lazy cleanup of stale entries during traversal.

#### **3.5. New Coarsening approaches in Nested Dissection**

#### **3.6. Hypergraph Based Ordering**

#### **3.7. GPU Implementation of RCM**

GPU Implementation of RCM is basically a parallel implementation of the breadth-first search (BFS) algorithm. Much research is available on parallel BFS, and the implementation in this thesis is based on the NVIDIA work on GPU-accelerated BFS in [7].

This GPU breadth-first search implementation uses a level-synchronous algorithm that processes the graph one depth level at a time, maintaining two key data structures: a vertex frontier (vertices to be explored in the current iteration) and an edge frontier (all neighbors of vertices in the current frontier). The algorithm begins with a single source vertex and alternates between two fundamental operations across BFS levels.

### 3. Implementation and Optimizations

---

**Algorithm 8:** Parallel AMD Algorithm

---

```
input :Matrix  $A$  ( $n \times n$ ), parameters mult, lim
output:Elimination ordering

1 Preprocessing;;
2  $G \leftarrow \text{initialize\_quotient\_graph}(A + A^T)$  // Compute symmetric pattern
3 Initialize degree lists for each thread;;
4 forall  $tid = 0$  to  $\text{num\_threads} - 1$  in parallel do  $\text{initialize\_degree\_list}(tid)$  ;
5 Main elimination loop;;
6 while  $|V| > 0$  do
7   Find minimum approximate degree across all threads;;
8    $\text{amd} \leftarrow \text{find\_global\_minimum\_degree}()$ ;
9   Select distance-2 independent set of pivots;;
10   $D \leftarrow \text{distance\_2\_independent\_set}(\text{amd}, \text{mult}, \text{lim})$ ;
11  if  $|D| = 0$  then
12    | break // No more valid pivots
13  end
14  Eliminate pivots in parallel;;
15  forall pivot  $p \in D$  in parallel do  $\text{tid} \leftarrow \text{get\_thread\_id}()$ ;
16   $\text{eliminate\_pivot}(\text{tid}, p)$ ;
17  ;
18  Barrier synchronization;;
19  barrier();
20 end
21 return elimination ordering;
```

---

### *3. Implementation and Optimizations*

Each BFS iteration follows a two-phase process. In the expansion phase, the algorithm takes the current vertex frontier and performs parallel neighbor gathering to create the edge frontier. Multiple threads cooperatively read the adjacency lists of frontier vertices from the compressed sparse row (CSR) representation, collecting all outgoing edges. In the contraction phase, the algorithm filters this edge frontier by checking each neighbor's visitation status, removing already-visited vertices and duplicates to produce the vertex frontier for the next iteration. This process continues until the vertex frontier becomes empty, indicating the traversal is complete.

The expansion phase uses a multi-granularity approach to handle the irregular degree distributions common in real graphs. For vertices with small adjacency lists, the algorithm employs scan-based gathering where threads use prefix sum to compute scatter offsets, creating a perfectly packed array of neighbors that allows all threads to participate in memory reads without SIMD lane waste. For medium-sized adjacency lists, warp-based gathering assigns entire 32-thread warps to cooperatively process single vertices, with threads strip-mining through the adjacency list in parallel. For very large adjacency lists, CTA-based gathering enlists entire thread blocks (hundreds of threads) to process individual high-degree vertices. This hybrid strategy automatically adapts to the workload characteristics and ensures efficient GPU utilization regardless of degree distribution.

I took an already existing implementation of the parallel BFS algorithm from the NVIDIA CUDA SDK [8] and adapted it for reordering the graph using RCM.

# Chapter 4

## Results

In this chapter, we first describe our evaluation setup, which consists of the HPC cluster used for our experiments, the matrix datasets, and the algorithmic implementations evaluated. We then present a detailed analysis of various performance metrics, including execution time, memory usage, parallelization and matrix fill-in reduction.

### 4.1. Evaluation setup

Everything that is evaluated is ran on Fritz, a high-performance computing cluster at NHR@FAU.

#### 4.1.1. Hardware Infrastructure

Fritz is a parallel CPU cluster operated by NHR@FAU, featuring Intel Ice Lake and Sapphire Rapids processors with an InfiniBand (IB) network and a Lustre-based parallel filesystem accessible under \$FASTTMP. The cluster configuration consists of:

Nodes	CPUs and Cores	Memory	Slurm Partition
992	2 × Intel Xeon Platinum 8360Y (Ice Lake) 2 × 36 cores @2.4 GHz	256 GB	singlenode, multinode
48	2 × Intel Xeon Platinum 8470 (Sapphire Rapids) 2 × 52 cores @2.0 GHz	1 TB	spr1tb
16	2 × Intel Xeon Platinum 8470 (Sapphire Rapids) 2 × 52 cores @2.0 GHz	2 TB	spr2tb

Table 4.1.: Fritz cluster node configuration

#### 4. Results

The login nodes fritz[1-4] are equipped with  $2 \times$  Intel Xeon Platinum 8360Y (Ice Lake) processors and 512 GB main memory. Additionally, the remote visualization node fviz1 features  $2 \times$  Intel Xeon Platinum 8360Y processors with 1 TB main memory, one Nvidia A16 GPU, and 30 TB of local NVMe SSD storage.

#### 4.2. Matrices dataset

We selected a diverse set of sparse matrices from the SuiteSparse Matrix Collection and our custom dataset of matrices encountered in Quantum Transport to evaluate our implementations. The SuiteSparse matrices chosen are usually which were commonly used in the literature for *METIS* and *SCOTCH* own evaluations. [TODO about Quantum transport matrices]

Graph name	Group	No. of vertices	No. of edges
copter2	CFD	55476	407714
ex10	CFD	2410	28625
parabolic_fem	CFD	525825	2100225
3Dspectralwave	Chemistry	680943	17165766
CO	Chemistry	221119	3943588
crystk03	Chemistry	24696	887937
Ga19As19H42	Chemistry	133123	4508981
Si87H76	Chemistry	240369	5451000
SiH4	Chemistry	5041	88472
circuit_3	Circuit	12127	48137
G2_circuit	Circuit	150102	438388
memplus	Circuit	17758	126150
rajat06	Circuit	10922	28922
rajat10	Circuit	30202	80202
rajat15	Circuit	37261	443573
144	Graphs	144649	1074393
598a	Graphs	110971	741934
auto	Graphs	448695	3314611
ca-AstroPh	Graphs	18772	198110
citationCiteseer	Graphs	268495	1156647
delaunay_n20	Graphs	1048576	3145686
dictionary28	Graphs	52652	89038
boyd2	Optimizations	466316	890091
c-55	Optimizations	32780	218115
c-70	Optimizations	68924	363955
gupta2	Optimizations	62064	2155175
gupta3	Optimizations	16783	4670105
lpl1	Optimizations	32460	180248

*Continued on next page*

#### 4. Results

Graph name	Group	No. of vertices	No. of edges
ncvxbqp1	Optimizations	50000	199984
apache1	Structural	80800	311492
apache2	Structural	715176	2766523
nasasrb	Structural	54870	1366097

Table 4.2.: SuiteSparse matrix dataset used for evaluation

Table 4.3.: Quantum Transport matrix dataset used for evaluation

Graph name	No. of vertices	No. of edges	Description
airpoll_conditional_st3	8499	1218651	Spatio-temporal statistical Modelling
airpoll_prior_st3	8499	1205931	Spatio-temporal statistical Modelling
cnt_cp2k	9152	6154350	DFT Hamiltonian of Carbon Nanotube
cnt_w90	768	71680	DFT Hamiltonian of Carbon Nanotube
kmc_potential_0	403605	10007089	Kinetic Monte Carlo for Device Modelling
kmc_potential_1	1632355	41208963	Kinetic Monte Carlo for Device Modelling
qubit_fem_b4	116380	2517228	FEM Matrix for Qubit
sinw_w90	7488	5310160	DFT Hamiltonian of Silicon Nanowire

### 4.3. Methodology

For our evaluation, we systematically assessed the performance of eleven different reordering algorithms across all matrices listed in Tables 4.2 and 4.3. The evaluated algorithms encompass both sequential and parallel approaches, ranging from classical degree-based methods to graph partitioning and hypergraph-based techniques:

#### Classical Methods:

- **AMD** – Approximate Minimum Degree
- **COLAMD** – Column Approximate Minimum Degree
- **RCM** – Reverse Cuthill-McKee
- **Natural** – Original matrix ordering (baseline)

#### Graph Partitioning Methods:

- **METIS** – Sequential graph partitioning-based reordering
- **METIS+IC** – METIS with improved coarsening integration
- **SCOTCH** – Sequential graph partitioning
- **PT-SCOTCH** – Parallel graph partitioning
- **ParMETIS** – Parallel graph partitioning

## 4. Results

### Hypergraph-based Methods:

- **HG-2, HG-4, HG-8, HG-16** – Hypergraph-based reordering with block partition sizes of 2, 4, 8, and 16 vertices respectively

The hypergraph-based methods (HG-2, HG-4, HG-8, HG-16) employ different block partition sizes to explore the trade-off between computational complexity and reordering/parallelization quality. All these algorithms are described in detail in Chapter 2.

#### 4.3.1. Performance Metrics and Measurement Protocol

Our comprehensive evaluation measures five key performance indicators to assess the effectiveness of each reordering algorithm:

**Fill-in Reduction:** We measure the number of non-zero entries introduced during symbolic Cholesky factorization using the `cmpfillin` binary provided by the METIS library. This tool performs symbolic factorization (as described in Chapter 2) to count the additional non-zero entries without executing the actual numerical computation.

**Computational Complexity:** We track the total floating-point operation count required for factorization, providing insight into the computational efficiency gains achieved by different reordering strategies.

**Reordering Time:** The wall-clock time required to compute the reordering itself is measured using Python’s `time.perf_counter()` with microsecond precision. This overhead cost is crucial for understanding the practical applicability of each method.

**Memory Usage:** We implement a dedicated memory monitoring system using a separate thread that samples memory consumption every 50 milliseconds during reordering operations. The monitoring captures both peak memory usage (maximum RSS observed) and average memory consumption throughout the reordering process. Memory measurements are obtained via the `psutil` library, tracking the resident set size (RSS) of the process and computing the memory increase relative to a baseline measured before algorithm execution.

**Parallelization Potential:** We assess the elimination tree depth as an indicator of the inherent parallelism available in the factorization. As mentioned in Chapter 2, the elimination tree is constructed by identifying parent-child relationships in the symbolic factorization structure, where the tree depth represents the critical path length and thus the theoretical lower bound on parallel factorization time.

For memory profiling, we initially attempted to use the Massif profiler from Valgrind, but encountered issues with our Python-based implementations and the cluster environment. Consequently, we developed a custom memory measurement protocol that

#### 4. Results

employs a thread-safe monitoring system to capture accurate memory usage patterns during reordering operations.

Our memory measurement implementation works by first establishing a baseline memory consumption before starting any reordering algorithm, recording the initial resident set size (RSS) using `psutil.Process().memory_info().rss`. During algorithm execution, a dedicated monitoring thread runs concurrently, sampling memory usage every 50 milliseconds throughout the entire operation. This thread continuously updates the maximum observed memory consumption and maintains a comprehensive list of all samples for subsequent average computation.

Upon completion of the reordering algorithm, we compute both the peak memory increase (calculated as the maximum observed RSS minus the baseline) and the average memory increase (computed as the mean of all samples minus the baseline), with results reported in kilobytes. For cases where the monitoring thread fails to capture meaningful data, such as very fast operations that complete before sufficient samples can be collected, we provide conservative estimates based on matrix characteristics, using approximately 10% of the matrix storage size as an upper bound for memory consumption.

All experiments are executed on the Fritz cluster to ensure consistent hardware conditions across all measurements. Each algorithm is tested on every matrix in our dataset, and results are averaged over five independent runs to mitigate variability.

#### 4.4. Evaluation

For each matrix in our dataset with different reordering algorithms applied, we have collected all the aforementioned metrics, and these can be found in our appendix/supplementary material A. Here, we present results and visualizations for our underlying metrics for some of them.

Figure 4.1 presents a comparative analysis of fill-in reduction across different matrix categories, showing the ratio of fill-in to original non-zeros for representative matrices from each domain. This visualization demonstrates how reordering effectiveness varies significantly across different problem types and matrix structures.

## 4. Results

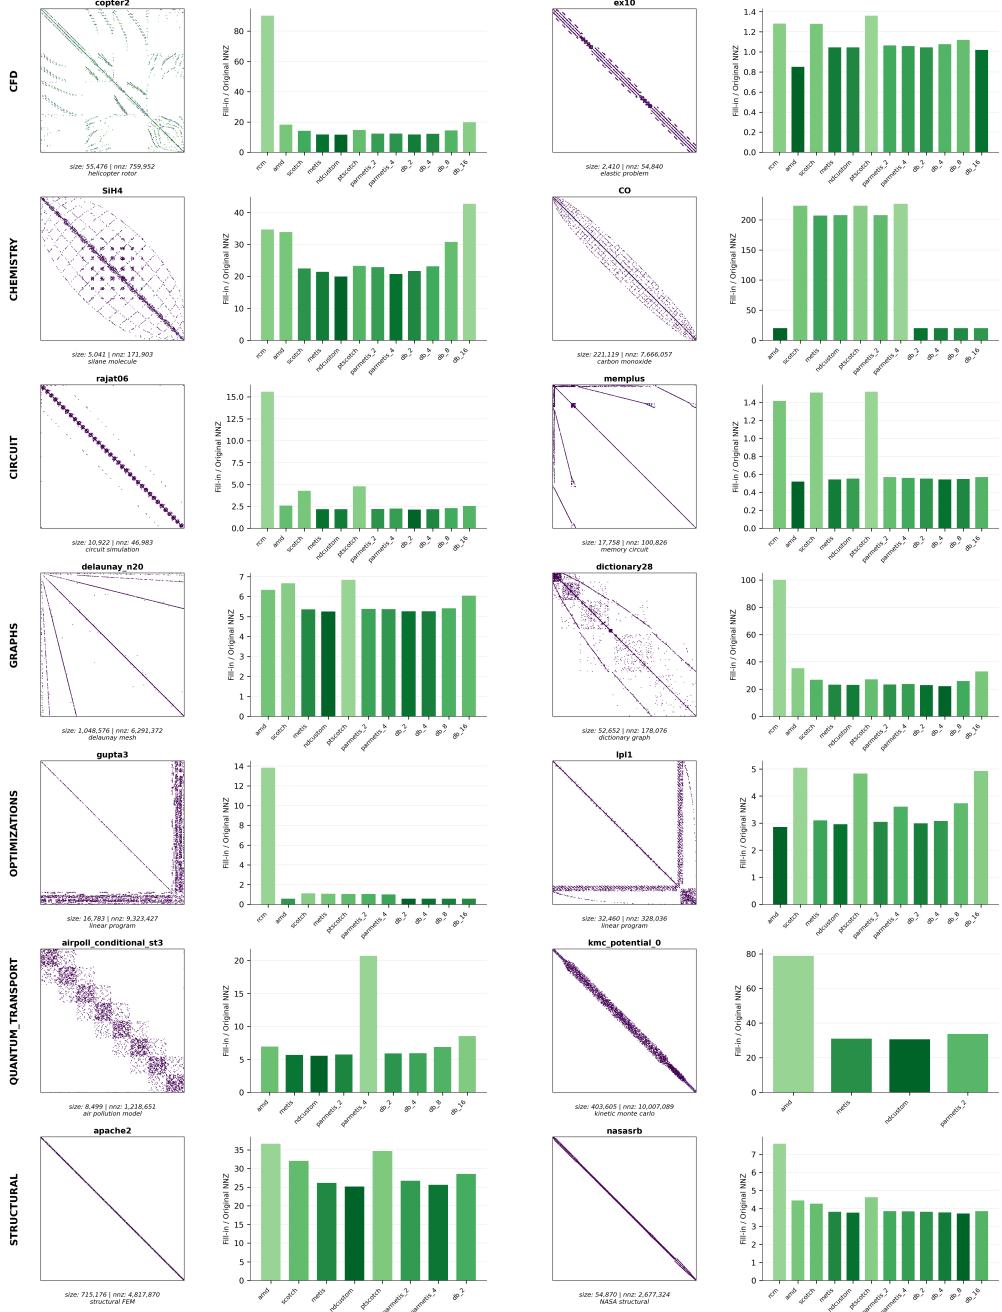


Figure 4.1.: Fill-in ratio (fill-in/original nnz) comparison across different matrix categories

Our experimental results reveal distinct performance patterns across different matrix domains and algorithm categories. The hypergraph-based methods, METIS, and our custom METIS implementation with improved coarsening (METIS+IC) consistently

#### 4. Results

emerge as the best-performing approaches, demonstrating better fill-in reduction compared to classical methods. These methods usually achieve 30-60% lower fill-in values than traditional approaches like AMD and COLAMD across most domains. Among the hypergraph methods, smaller partition sizes (HG-2 and HG-4) typically produce optimal results, while larger partitions are ineffective. Our custom METIS+IC implementation validates the effectiveness of enhanced coarsening strategies by achieving competitive performance with the best hypergraph methods. SCOTCH and PT-SCOTCH show moderate improvements over classical methods but do not match the quality of METIS-based approaches.

Domain-specific analysis reveals that hypergraph methods, METIS, and METIS+IC excel particularly in circuit and graph problems, where they achieve 35-50% fill-in reduction compared to AMD. For CFD matrices, these advanced methods maintain their superiority with modest but consistent improvements (10-30%), while chemistry matrices present scalability challenges where many classical methods fail on the largest instances, leaving primarily METIS-based or classical methods viable. The optimization and structural domains show variable performance depending on problem structure, with hypergraph methods and METIS variants providing 10-25% improvements in most cases.

Figure 4.2 presents the scaling behavior of reordering algorithms with respect to the number of non-zero entries. This graph shows the time required for each algorithm to complete and how it scales with increasing non-zero counts.

The reordering time results show notable performance differences across methods. RCM and AMD are the fastest options, with RCM typically completing in under one second (0.006-0.76s) across all matrix sizes and AMD performing well for small-to-medium matrices (0.046-2.09s). SCOTCH and METIS generally complete within reasonable timeframes (under 50 seconds), while parallel methods like ParMETIS and PT-SCOTCH have higher overhead that may be justified for larger problems. The hypergraph partitioning methods (HG variants) require significantly more time, with some executions taking over 270 seconds, which limits their practical applicability for interactive use.

Regarding scalability with increasing NNZ, the methods exhibit different growth patterns. AMD shows sublinear scaling (slope 0.52), which becomes more favorable as matrix size increases, while RCM maintains near-linear scaling (slope 0.72) with consistent performance characteristics. SCOTCH demonstrates reasonable scaling (slope 0.65), and the ParMETIS variants maintain moderate complexity (slopes 0.64-0.68) despite their higher base execution times. METIS+IC shows superlinear scaling (slope 1.14) due to its modified coarsening approach, which becomes problematic for very large matrices, while METIS approaches linear scaling (slope 0.81). The HG hypergraph methods show scaling slopes between 0.71-0.76, though their high absolute execution times remain a concern.

The results indicate that simpler methods like AMD and RCM often provide better practical performance than more complex approaches for large-scale sparse matrix reordering. While parallel methods may become advantageous for extremely large problems, the hypergraph partitioning approaches appear computationally expensive relative to their benefits in typical sparse matrix applications.

## 4. Results

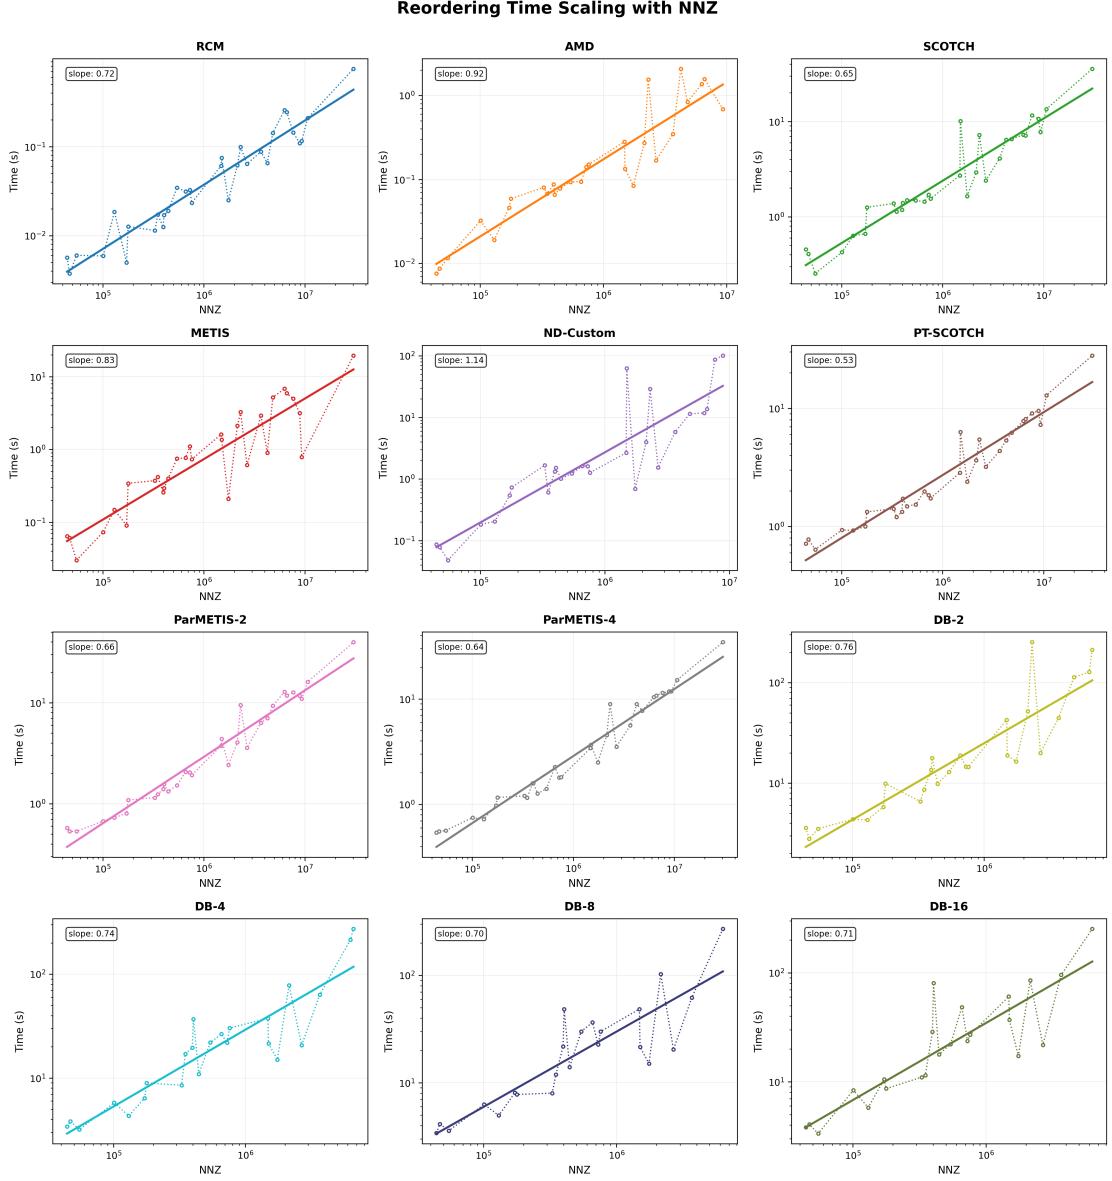


Figure 4.2.: Reordering time scaling with matrix size (number of non-zeros) for different algorithms. Each subplot shows the relationship between matrix size and computational time for various reordering methods.

We then now look at the elimination tree depth as a measure of parallelism. Similar to fill-in, we sample randomly two matrices from different domains to show the elimination tree depth for different reordering algorithms.

The elimination tree depth results provide insight into the parallelization potential of different reordering methods, as shallower trees enable more parallel execution

#### 4. Results

during factorization. The parallel methods show clear advantages in reducing tree depth compared to sequential approaches. ParMETIS variants consistently produce the shallowest elimination trees, with ParMETIS-2 and ParMETIS-4 often achieving depths in the range of 14-42 across various matrices, while ParMETIS-8 and ParMETIS-16 typically maintain depths between 25-50. METIS also performs well for parallelization potential, frequently producing trees with depths around 22-160, which represents a significant improvement over many sequential methods.

Sequential methods show mixed results for parallelization potential. AMD typically produces tree depths ranging from very shallow (4-20) for some matrices to quite deep (250-830) for others, making its parallelization benefits matrix-dependent. NDCUSTOM, being METIS with modified coarsening, generally produces similar tree depths to METIS, typically ranging from 24-171, though it can occasionally produce deeper trees (up to 839) for some matrices, suggesting that the coarsening modification does not consistently improve parallelization potential. RCM consistently produces the deepest trees, often exceeding 300-400 in depth and reaching over 700 for some matrices, which severely limits parallel factorization opportunities. SCOTCH performs moderately, with tree depths generally ranging from 23-290, while PT-SCOTCH shows improved parallelization potential with depths typically between 54-328, though it occasionally produces very deep trees exceeding 700.

The hypergraph partitioning methods (DB variants) show variable performance for parallelization, with tree depths ranging from 19-301 depending on the number of partitions used. DB-2 often produces deeper trees (49-603), while DB-16 tends toward moderate depths (66-830), though the relationship between partition count and tree depth is not always consistent. Natural ordering predictably produces very deep trees (often 200-800+), confirming the necessity of reordering for parallel factorization. The results indicate that while parallel reordering methods require more computational time, they provide substantial benefits for subsequent parallel factorization phases through significantly reduced elimination tree depths.

## 4. Results

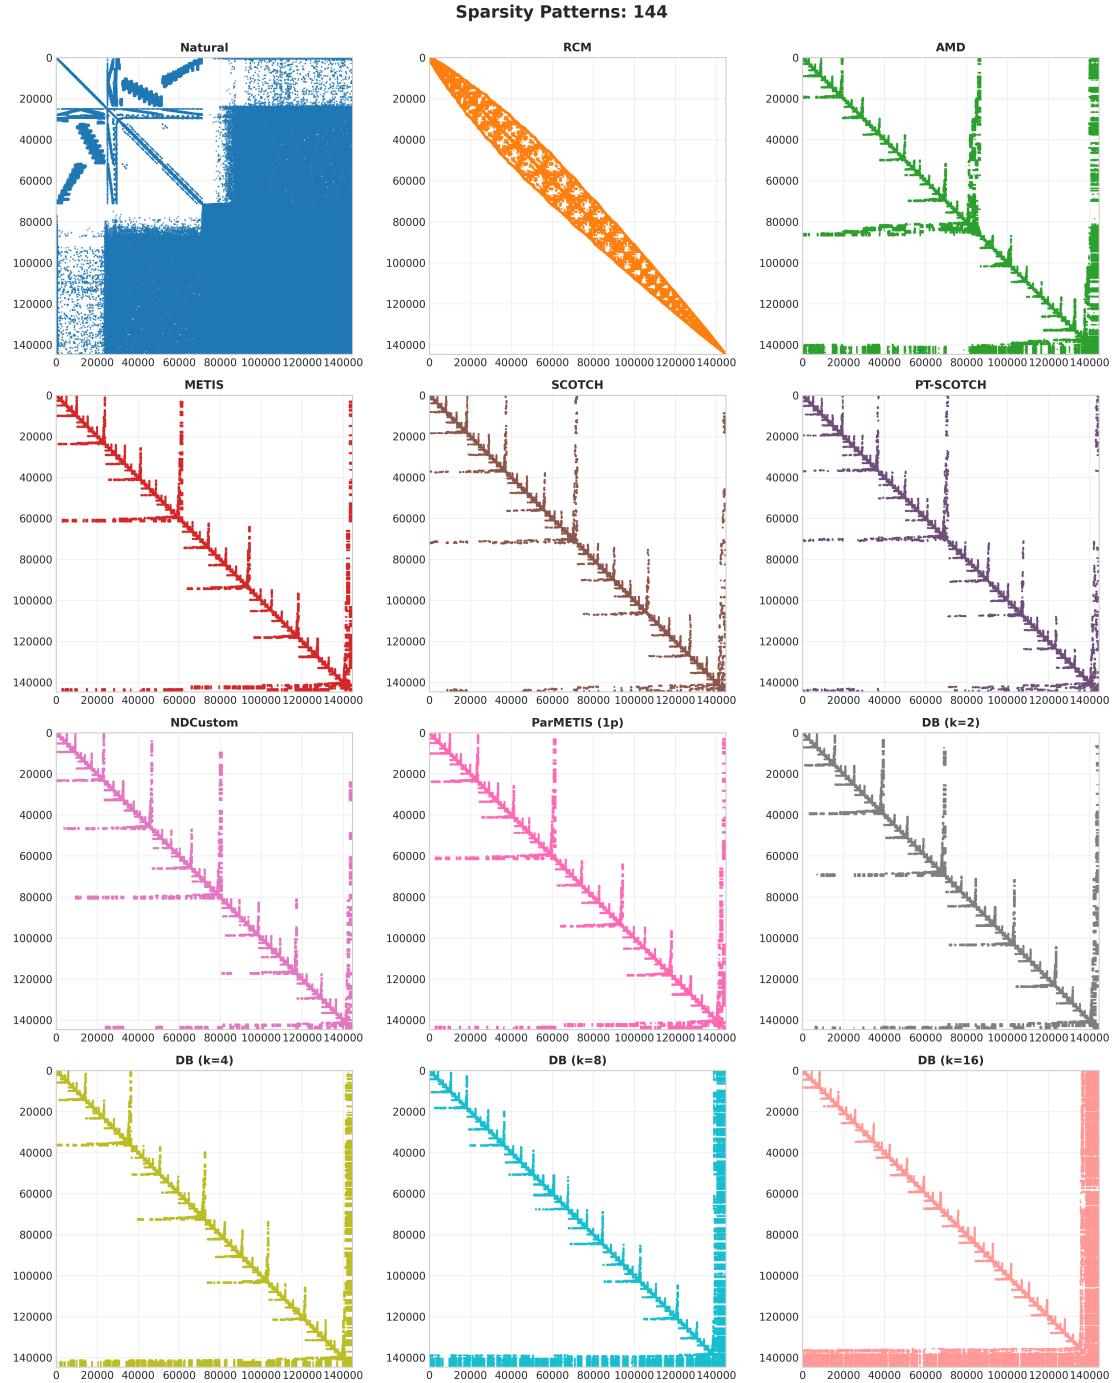


Figure 4.3.: Sparsity patterns of matrix 144 with different reordering algorithms applied. The natural ordering (top-left) shows a scattered pattern leading to high fill-in, while optimized reorderings (other panels) demonstrate more defined structures.

## 4. Results

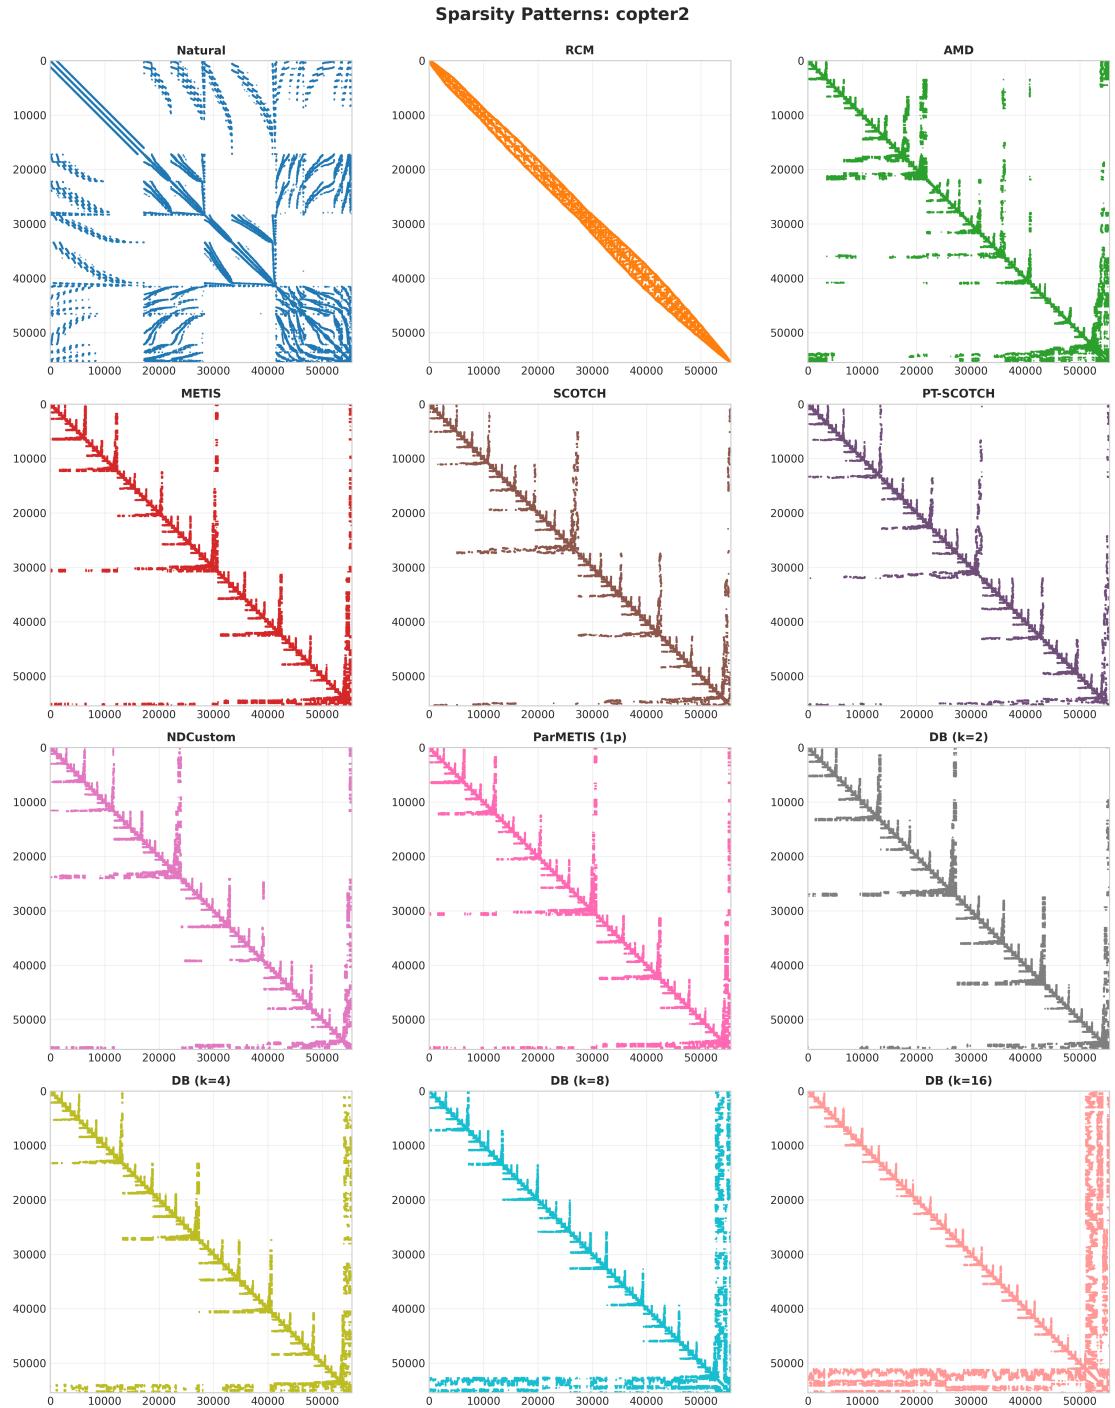


Figure 4.4.: Similarly, sparsity patterns of matrix copter2 with different reordering algorithms applied.

## 4. Results

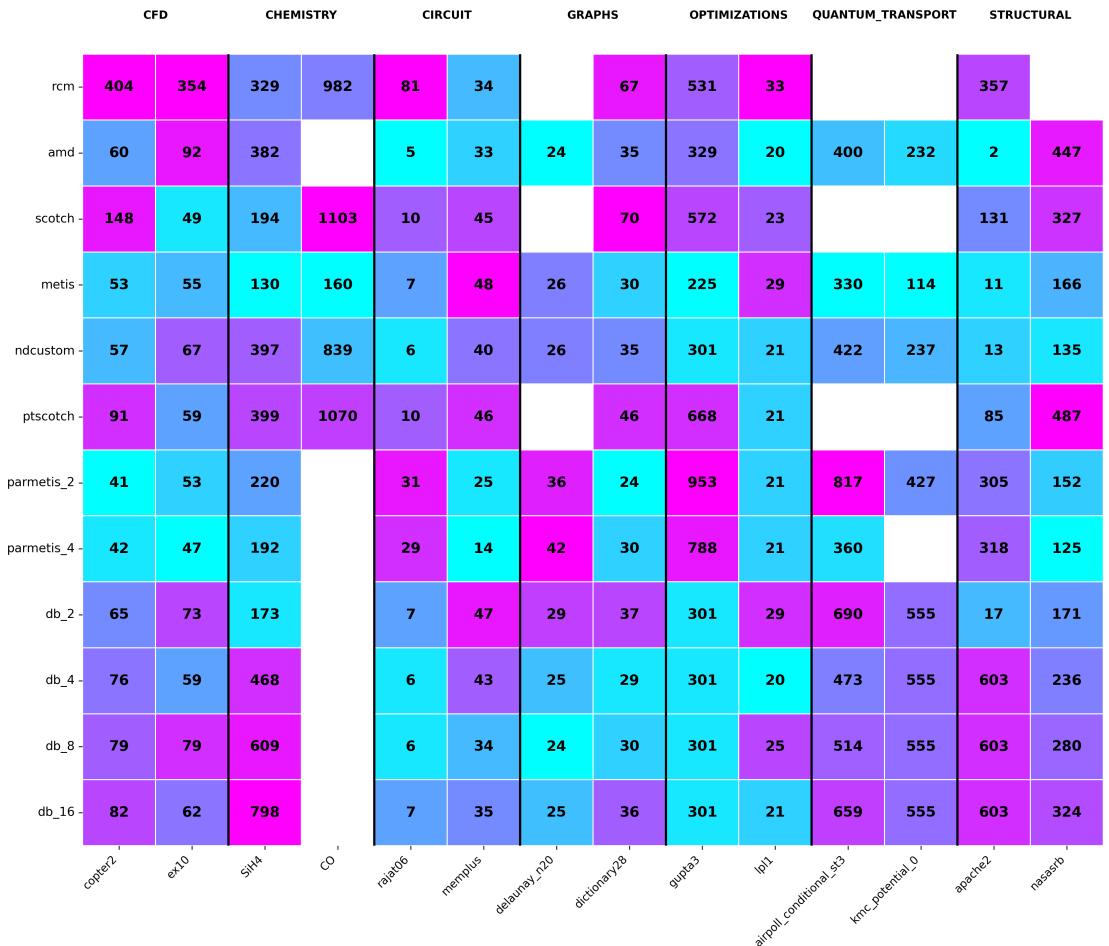


Figure 4.5.: Elimination tree depth comparison across different matrix categories for various reordering algorithms. Lower depth indicates better parallelization potential. Blue indicates lower depth (better), while pink indicates higher depth (worse).

## 4. Results

### GPU-accelerated RCM

In this section, we present the performance of our GPU-accelerated implementation of the Reverse Cuthill-McKee (RCM) algorithm compared to a traditional CPU-based implementation. We measure only the runtime of the RCM algorithm and not the fill-in, as the quality of the ordering produced by both implementations is identical.

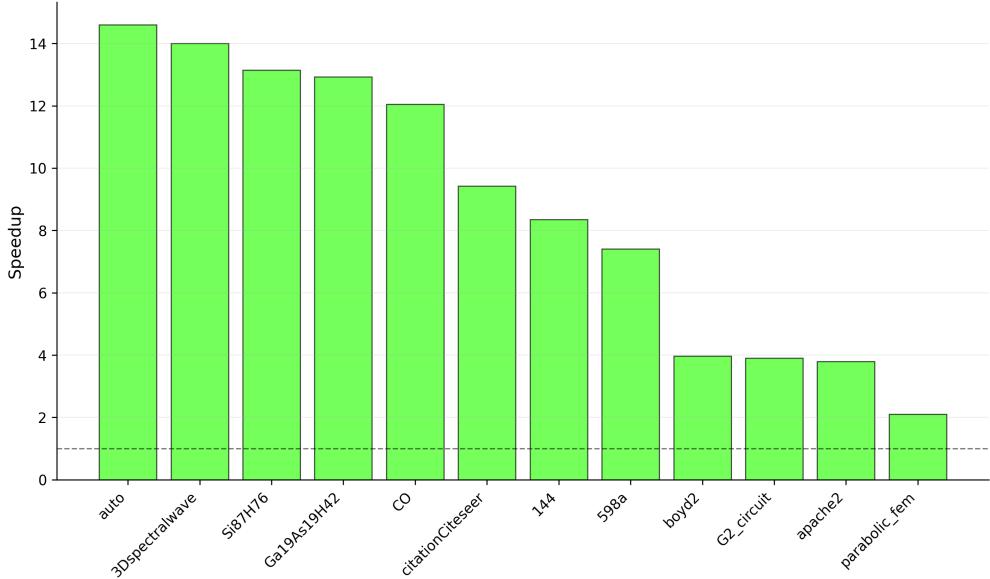


Figure 4.6.: Speedup of GPU-accelerated RCM compared to CPU-based RCM on large matrices (over 10K nodes).

### 4.5. Current limitations

Our evaluation has important limitations that should be acknowledged. First, our analysis is not exhaustive—we evaluated only a subset of available sparse matrix reordering algorithms and libraries, focusing on the most commonly used and representative methods. Additionally, our memory measurement system faced significant challenges: the monitoring thread frequently failed to capture meaningful data for fast-completing algorithms or when external binaries changed process IDs (particularly with hypergraph methods). For several large matrices, symbolic factorization failed due to memory constraints, preventing complete evaluation of all methods across our entire dataset.

The practical applicability of some methods is also limited by performance constraints. Hypergraph-based methods, while sometimes producing high-quality reorderings, proved computationally expensive and impractical for large matrices. Furthermore, our GPU acceleration efforts were limited to the RCM algorithm—we did not implement GPU versions of other reordering methods, which represents a significant opportunity for future performance improvements across the entire algorithm suite.

# Chapter 5

## Other Approaches

In this chapter, we explore alternative and promising approaches to the sparse matrix reordering problem beyond traditional heuristic methods. These methods were found by myself to be either computationally difficult to scale for large matrices or were found to have some bottlenecks, but nevertheless represent interesting directions for future work.

### 5.1. Graph Reinforcement Learning for Reordering

As we know that sparse symmetric matrices can be represented as undirected graphs, it seems promising that this representation allows GNNs to naturally capture the local neighborhood relationships that classical ordering heuristics rely on, such as node degree and clustering patterns.

Traditional heuristics like minimum degree ordering make greedy decisions based on limited local information. There has been a few approaches for using machine learning to learn better heuristics, such as the work by [5], which uses CNN and they tackle large matrices by partitioning the graph.

GNNs, however, can propagate information across multiple hops in the graph, enabling each node to consider not just its immediate neighbors but also the broader structural context. It also allows us to have variable sized graphs. This multi-hop reasoning capability allows the network to anticipate how eliminating one node will affect distant parts of the matrix, potentially leading to more effective ordering decisions.

The message-passing architecture of GNNs naturally models the fill-in process during matrix factorization. When a node is eliminated, it creates new connections between its neighbors, which GNNs can represent through their aggregation and update mechanisms. The network can learn to predict these fill-in patterns and make elimination choices that minimize overall structural complexity.

The model shows promising results on small matrices, as shown in Table 5.1, but struggles to scale to larger matrices due to the computational complexity of training

## 5. Other Approaches

Matrix	vs RCM (%)	vs MinDeg (%)
ash85	9.8	81.4
bcsstk22	28.4	51.0

Table 5.1.: GNN-based method performance comparison with RCM and MinDegree

GNNs on large graphs. Future work can explore scaling or coarser representations of the graph to make it feasible for larger matrices.

## 5.2. GPU Accelerated Nested Dissection

There is very limited work on GPU-accelerated reordering algorithms. I had presented earlier in this thesis, a GPU implementation of the RCM algorithm. Another promising approach is to implement the nested dissection algorithm on GPUs. One such approach is presented by [9], which leverages a GPU-based geometry processing library to perform fast graph partitioning, a key step in nested dissection. The method achieves significant speedups over the traditional CPU-based METIS library but the quality of the ordering is around 3 to 8 times worse than METIS, in their experiments.

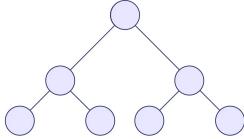


Figure 5.1.: Recursive calls in nested dissection

The primary bottleneck in running nested dissection on GPUs is the recursive nature of the algorithm. Each level of recursion requires partitioning the graph and identifying separators. We leveraged a graph processing library, Jet, which is a multilevel graph partitioning library designed for GPU [10]. The library comes with an efficient implementation of k-way partitioning done using a parallel multilevel approach. One can do a 2-way partitioning by setting  $k=2$ , but unsurprisingly, it performs worse in run-time as each call/partition (circle in Figure 5.1) is done sequentially. Another approach was to use  $k = 2^d$ , where  $d$  is the depth of recursion such each subgraph reaches the threshold size after  $d$  levels of recursion.

This approach is similar to the approach done using Hypergraph partitioning as mentioned in Chapter 3, in the way that the partition divides the graph into smaller subgraphs and the cumulative separators found are placed at the end, producing sparsity-pattern similar to hypergraph based methods in Chapter 3. This method, which is questionable to call it Nested Dissection, however produces orderings faster than METIS for large graphs but the quality of the ordering is still not comparable to METIS expectedly as you get better separators when you do recursive bisection than a single k-way partitioning. Such an approach produces orderings that are around 2 to 3 times worse than METIS in my experiments on large matrices (`parabolic_fem.graph` and `auto.graph`) but the speedups are significant around 5 to 6 times tested on the same environment as mentioned in Chapter 4.

# Chapter 6

## Conclusion and Future Work

This thesis presented a comprehensive evaluation of sparse symmetric matrix reordering algorithms across diverse matrix structures and application domains. Through systematic benchmarking on the Fritz HPC cluster, we assessed many reordering algorithms, ranging from classical degree-based methods to newer graph partitioning and hypergraph-based techniques, across multiple performance metrics including fill-in reduction, computational efficiency, memory consumption, and parallelization potential.

Our evaluation reveals that advanced methods such as METIS, SCOTCH, and hypergraph based approaches consistently outperform classical algorithms like AMD and RCM in terms of fill-in reduction, achieving 30-60% fewer fill-ins across most domains. Among the hypergraph methods, smaller partition sizes (HG-2 and HG-4) demonstrate superior effectiveness. Our custom METIS implementation with improved coarsening (METIS+IC) validates the benefits of enhanced coarsening strategies, producing results competitive with the best hypergraph methods. However, this improvement comes at a computational cost. Hypergraph methods require significantly more time (up to 270 seconds), while simpler methods like RCM complete in under one second.

In the course of this work, we implemented several efficient variants of existing algorithms. Our GPU-accelerated RCM implementation demonstrates substantial speedups over CPU-based versions for large matrices. The METIS+IC variant, incorporating refined coarsening strategies, achieves fill-in quality comparable to more complex hypergraph methods while maintaining reasonable execution times. We also explored hypergraph-based reordering methods that exploit block structures in matrices, showing promise particularly for circuit and graph problems.

Several promising directions remain for future investigation. Machine learning approaches, particularly graph neural networks (GNNs), show encouraging results on small matrices but require techniques like coarsening or hierarchical representations to scale to larger problems. Reordering algorithms for emerging architectures also warrant further research, while we implemented GPU-accelerated RCM, extending GPU implementations to nested dissection, minimum degree variants, and hypergraph methods could yield significant performance improvements. Finally, our evaluation re-

## *6. Conclusion and Future Work*

veals that algorithm effectiveness varies significantly across matrix domains, suggesting that adaptive schemes which automatically select or tune algorithms based on matrix characteristics could improve performance across diverse applications.

Appendix **A**

# Experimental Results

This appendix presents the complete experimental results for all matrix reordering algorithms evaluated in this thesis.

Table A.1 presents the number of nonzero entries after factorization for each reordering method. Table A.2 presents the computational time in seconds required to compute each reordering. Lower values indicate faster reordering algorithms. Table A.3 presents the peak memory usage in kilobytes during the reordering computation. Lower values indicate better memory efficiency. Table A.4 presents the average memory usage in kilobytes during the reordering computation. Table A.5 presents the depth of the elimination tree produced by each reordering. Lower depths generally indicate better potential for parallel factorization.

Empty or zero cells indicate that the algorithm either failed to complete, exceeded memory limits, timed out or some computation error. All measurements were performed on identical hardware configurations to ensure fair comparison. There are some meaningless values for memory usage, especially for hypergraph methods, for example showing as 4KB, very likely caused by failing the capture the memory usage for the corresponding PID process.

## A. Experimental Results

Matrix	Natural	RCM	AMD	METIS	NDCustom	ParMETIS-2	ParMETIS-4	ParMETIS-8	SCOTCH	PT-SCOTCH	db-2	db-4	db-8	db-16
144	$1.8 \times 10^9$	$9.3 \times 10^8$	$4.7 \times 10^7$	$4.6 \times 10^7$	$4.7 \times 10^7$	$4.8 \times 10^7$	$5.1 \times 10^7$	$5.3 \times 10^7$	$4.7 \times 10^7$	$4.7 \times 10^7$	$5.2 \times 10^7$	$6.8 \times 10^7$	$6.8 \times 10^7$	$6.8 \times 10^7$
3Dspectralwave	$1.8 \times 10^9$	—	$1.3 \times 10^9$	$1.3 \times 10^9$	$1.4 \times 10^9$	$1.4 \times 10^9$	$2.6 \times 10^7$	$2.6 \times 10^7$	$1.4 \times 10^9$	$2.0 \times 10^9$	$1.8 \times 10^9$	$1.8 \times 10^9$	$1.8 \times 10^9$	$1.8 \times 10^9$
598a	$1.6 \times 10^8$	$4.6 \times 10^7$	$4.6 \times 10^7$	$2.5 \times 10^7$	$2.5 \times 10^7$	$2.5 \times 10^7$	—	$2.6 \times 10^7$	$2.9 \times 10^7$	$2.6 \times 10^7$	$2.6 \times 10^7$	$2.9 \times 10^7$	$2.9 \times 10^7$	$2.9 \times 10^7$
CO	$1.6 \times 10^9$	$1.8 \times 10^{19}$	$1.6 \times 10^8$	$1.6 \times 10^8$	$1.6 \times 10^9$	$1.6 \times 10^9$	$1.6 \times 10^9$	$1.7 \times 10^9$	$1.7 \times 10^9$	$1.6 \times 10^8$				
G2_circuit	$1.0 \times 10^9$	$1.2 \times 10^8$	$9.9 \times 10^6$	$6.4 \times 10^6$	$6.4 \times 10^6$	$6.4 \times 10^6$	$6.6 \times 10^6$	$6.7 \times 10^6$	$7.8 \times 10^6$	$8.0 \times 10^8$	$7.9 \times 10^8$	—	—	$9.0 \times 10^6$
Ga19_Asi19H42	—	—	—	$7.5 \times 10^8$	$7.3 \times 10^8$	$7.7 \times 10^8$	$1.9 \times 10^9$	$2.0 \times 10^9$	$2.0 \times 10^9$	—	—	—	—	—
Si87H76	$6.8 \times 10^6$	$5.8 \times 10^6$	$5.8 \times 10^6$	$3.7 \times 10^6$	$3.4 \times 10^6$	$3.7 \times 10^6$	$3.9 \times 10^6$	$3.9 \times 10^6$	$4.0 \times 10^6$	$3.7 \times 10^6$	$4.0 \times 10^6$	$5.3 \times 10^6$	$7.4 \times 10^6$	$7.4 \times 10^6$
SiH4	$6.5 \times 10^7$	$4.4 \times 10^7$	$1.3 \times 10^7$	$1.0 \times 10^7$	$9.9 \times 10^6$	$1.0 \times 10^7$	$1.0 \times 10^7$	$1.3 \times 10^7$	$1.3 \times 10^7$	$1.0 \times 10^7$	$1.0 \times 10^7$	$9.5 \times 10^6$	$1.1 \times 10^7$	$1.1 \times 10^7$
apache1	—	—	$1.3 \times 10^8$	$1.2 \times 10^8$	$1.2 \times 10^8$	$1.3 \times 10^8$	$1.3 \times 10^8$	$1.4 \times 10^8$	$1.4 \times 10^8$	$1.5 \times 10^8$	$1.5 \times 10^8$	$1.7 \times 10^8$	—	—
apache2	—	—	$5.7 \times 10^8$	$2.2 \times 10^8$	$2.4 \times 10^8$	$2.4 \times 10^8$	$2.4 \times 10^8$	$2.4 \times 10^8$	$2.3 \times 10^8$	$8.0 \times 10^5$				
auto	$2.7 \times 10^8$	$2.8 \times 10^8$	$8.1 \times 10^5$	$8.0 \times 10^5$	$8.9 \times 10^5$	$8.9 \times 10^5$	$8.9 \times 10^5$	$8.9 \times 10^5$	$1.6 \times 10^6$	$1.8 \times 10^9$	$8.0 \times 10^5$	$8.0 \times 10^5$	$8.0 \times 10^5$	$8.0 \times 10^5$
boyd2	$c-55$	—	$4.6 \times 10^6$	$3.5 \times 10^6$	$3.2 \times 10^6$	$3.4 \times 10^6$	$3.1 \times 10^6$	$3.1 \times 10^6$	$3.1 \times 10^6$	$3.0 \times 10^6$				
c-70	$3.0 \times 10^8$	$1.1 \times 10^8$	$4.0 \times 10^6$	$3.1 \times 10^6$	$3.1 \times 10^6$	$3.1 \times 10^6$	$3.2 \times 10^6$							
ca_AstroPh	—	—	$7.2 \times 10^9$	$8.0 \times 10^6$	$8.2 \times 10^6$	$8.2 \times 10^6$	$7.8 \times 10^6$	$1.0 \times 10^7$	$1.0 \times 10^7$	$7.0 \times 10^6$	$7.0 \times 10^6$	$9.1 \times 10^6$	$1.1 \times 10^7$	$1.1 \times 10^7$
circuit_3	$2.7 \times 10^4$	$1.3 \times 10^6$	$2.6 \times 10^4$	$2.9 \times 10^4$	$2.9 \times 10^4$	$2.9 \times 10^4$	$3.0 \times 10^4$	$2.9 \times 10^4$	$4.6 \times 10^4$	$4.6 \times 10^4$	$4.7 \times 10^4$	$2.8 \times 10^4$	$2.9 \times 10^4$	$2.9 \times 10^4$
citationCleseer	—	—	$6.7 \times 10^8$	$1.4 \times 10^7$	$8.9 \times 10^6$	$8.9 \times 10^6$	$8.9 \times 10^6$	$9.4 \times 10^6$	$9.4 \times 10^6$	$1.1 \times 10^7$	$1.1 \times 10^7$	$8.9 \times 10^6$	$9.3 \times 10^6$	$1.1 \times 10^7$
copier2	$7.0 \times 10^8$	$6.9 \times 10^7$	$1.4 \times 10^7$	$9.2 \times 10^6$	$1.0 \times 10^7$	$1.0 \times 10^7$	$9.2 \times 10^6$	$9.2 \times 10^6$	$9.2 \times 10^6$	$9.6 \times 10^6$				
crystk03	$1.4 \times 10^7$	$1.3 \times 10^7$	$1.0 \times 10^7$	$3.4 \times 10^7$	$3.3 \times 10^7$	$3.3 \times 10^7$	$3.4 \times 10^7$	$3.4 \times 10^7$	$4.2 \times 10^7$	$4.3 \times 10^7$	$3.3 \times 10^7$	$3.3 \times 10^7$	$3.4 \times 10^7$	$3.8 \times 10^7$
delanauy_n20	—	—	$4.0 \times 10^7$	$4.0 \times 10^7$	$9.4 \times 10^6$	$9.4 \times 10^6$	$9.7 \times 10^5$	$1.0 \times 10^6$	$1.0 \times 10^6$	$1.6 \times 10^6$	$9.8 \times 10^5$	$1.0 \times 10^6$	$1.2 \times 10^6$	$1.6 \times 10^6$
dictionary28	$7.9 \times 10^7$	$1.8 \times 10^7$	$6.3 \times 10^6$	$4.1 \times 10^6$	$4.1 \times 10^6$	$4.1 \times 10^6$	$4.3 \times 10^6$	$4.3 \times 10^6$	$4.8 \times 10^6$					
ex10	$9.1 \times 10^4$	$7.0 \times 10^4$	$4.7 \times 10^4$	$5.7 \times 10^4$	$5.7 \times 10^4$	$5.7 \times 10^4$	$5.7 \times 10^4$	$5.8 \times 10^4$	$5.8 \times 10^4$	$7.0 \times 10^4$	$7.0 \times 10^4$	$7.5 \times 10^4$	$7.5 \times 10^4$	$5.6 \times 10^4$
gupta2	$5.4 \times 10^8$	$4.7 \times 10^8$	$6.0 \times 10^6$	$5.8 \times 10^6$	$5.8 \times 10^6$	$5.8 \times 10^6$	$5.7 \times 10^6$	$5.7 \times 10^6$	$6.2 \times 10^6$	$7.4 \times 10^6$	$7.6 \times 10^6$	$5.4 \times 10^8$	$5.4 \times 10^8$	$5.4 \times 10^8$
gupta3	$5.4 \times 10^6$	$1.3 \times 10^8$	$5.4 \times 10^6$	$1.0 \times 10^7$	$8.6 \times 10^6$	$9.2 \times 10^6$	$9.9 \times 10^6$	$5.4 \times 10^6$	$5.4 \times 10^6$	$5.4 \times 10^6$				
lpl	—	—	$9.4 \times 10^5$	$1.0 \times 10^6$	$9.7 \times 10^5$	$9.7 \times 10^5$	$1.0 \times 10^6$	$1.0 \times 10^6$	$1.7 \times 10^6$	$1.6 \times 10^6$	$9.8 \times 10^5$	$1.0 \times 10^6$	$1.2 \times 10^6$	$1.6 \times 10^6$
memplus	$1.4 \times 10^5$	$5.2 \times 10^4$	$5.5 \times 10^4$	$5.6 \times 10^4$	$5.6 \times 10^4$	$5.6 \times 10^4$	$5.6 \times 10^4$	$5.7 \times 10^4$	$1.5 \times 10^5$	$1.5 \times 10^5$	$1.5 \times 10^5$	$5.6 \times 10^4$	$5.5 \times 10^4$	$5.7 \times 10^4$
nasasrb	$1.7 \times 10^7$	$2.0 \times 10^7$	$1.2 \times 10^7$	$1.0 \times 10^7$										
nevxbqp1	—	—	$4.2 \times 10^6$	$1.9 \times 10^6$	$1.9 \times 10^6$	$1.9 \times 10^6$	$2.5 \times 10^7$	$2.5 \times 10^7$	$2.5 \times 10^7$	$2.9 \times 10^7$	$2.9 \times 10^7$	$2.5 \times 10^7$	$2.6 \times 10^7$	$3.1 \times 10^7$
parabolic_fem	—	—	$3.6 \times 10^7$	$1.2 \times 10^5$	$1.0 \times 10^5$	$1.0 \times 10^5$	$1.0 \times 10^5$	$1.0 \times 10^5$	$2.0 \times 10^5$	$2.3 \times 10^5$	$1.0 \times 10^5$	$1.0 \times 10^5$	$1.1 \times 10^5$	$1.2 \times 10^5$
raja06	$2.0 \times 10^6$	$7.3 \times 10^5$	$3.3 \times 10^6$	$4.0 \times 10^5$	$3.4 \times 10^5$	$3.4 \times 10^5$	$3.4 \times 10^5$	$3.4 \times 10^5$	$6.3 \times 10^5$	$6.6 \times 10^5$	$6.6 \times 10^5$	$3.4 \times 10^5$	$3.5 \times 10^5$	$3.6 \times 10^5$
raja10	$7.4 \times 10^6$	$3.3 \times 10^6$	$4.0 \times 10^5$	$7.6 \times 10^5$	$7.7 \times 10^5$	$7.6 \times 10^5$	$7.7 \times 10^5$	$7.7 \times 10^5$	$1.0 \times 10^6$	$1.1 \times 10^6$	$7.9 \times 10^5$	$7.8 \times 10^5$	$8.5 \times 10^5$	$1.1 \times 10^6$
raja15	$8.1 \times 10^6$	—	$7.5 \times 10^7$	$6.8 \times 10^6$	$7.2 \times 10^6$	$7.2 \times 10^6$	$8.4 \times 10^6$	$7.0 \times 10^6$	$6.8 \times 10^6$	$7.0 \times 10^6$	$6.8 \times 10^6$	$7.3 \times 10^6$	$6.9 \times 10^6$	$1.0 \times 10^7$
airpoll_conditional_st3	$1.2 \times 10^7$	$2.5 \times 10^7$	$9.0 \times 10^6$	$8.4 \times 10^6$	$6.8 \times 10^6$	$7.1 \times 10^6$	$8.3 \times 10^6$	$7.1 \times 10^6$	$6.0 \times 10^6$	$6.0 \times 10^6$	$6.0 \times 10^6$	$6.8 \times 10^6$	$7.1 \times 10^6$	$9.9 \times 10^6$
airpoll_prior_st3	$6.0 \times 10^6$	$5.9 \times 10^6$	$8.4 \times 10^6$	$6.0 \times 10^6$	$6.0 \times 10^6$	$6.0 \times 10^6$	$5.3 \times 10^4$	$5.0 \times 10^4$	$3.5 \times 10^4$	$5.2 \times 10^4$	$1.0 \times 10^7$	$1.0 \times 10^7$	$1.0 \times 10^7$	$1.0 \times 10^7$
ent_cp2k	$3.5 \times 10^4$	$4.3 \times 10^4$	$7.9 \times 10^8$	—	—	—	—	$3.1 \times 10^8$	$1.1 \times 10^8$	$1.1 \times 10^8$	$1.1 \times 10^8$	$3.2 \times 10^8$	$3.4 \times 10^8$	—
cnt_w90	$2.7 \times 10^6$	$3.2 \times 10^8$	$3.8 \times 10^8$	$1.2 \times 10^8$	$1.1 \times 10^8$	$2.7 \times 10^6$	$5.0 \times 10^6$	$2.7 \times 10^6$	$5.0 \times 10^6$	$5.0 \times 10^6$	$4.9 \times 10^6$	$4.9 \times 10^6$	$5.0 \times 10^6$	$4.7 \times 10^6$
kmc_potential_0	$3.8 \times 10^8$	$2.7 \times 10^6$	$3.2 \times 10^8$	$2.7 \times 10^6$										
qubit_fem_b4	$2.7 \times 10^6$	$3.8 \times 10^8$	$3.2 \times 10^8$	$2.7 \times 10^6$										
sinrw_w90	—	—	—	—	—	—	—	—	—	—	—	—	—	—

Table A.1.: Fill-in (number of nonzeros after factorization) for each matrix and reordering algorithm.

## A. Experimental Results

Matrix	RCM	AMD	METIS	NDCustom	ParMETIS-2	ParMETIS-4	ParMETIS-8	SCOTCH	PT-SCOTCH	db-2	db-4	db-8	db-16
144	0.061	0.271	2.112	3.964	4.053	4.548	3.372	2.923	3.635	52.075	78.292	102.635	85.395
3Dspectralwave	0.746	0.0	19.562	0.0	39.755	34.463	30.827	35.659	27.972	0.0	0.0	0.0	0.0
598a	0.061	0.282	1.614	2.668	3.754	3.420	3.308	2.712	2.854	42.707	37.654	48.547	60.457
CO	0.144	0.0	4.985	86.865	12.658	11.394	10.552	11.558	9.062	0.0	0.0	0.0	0.0
G2_circuit	0.033	0.141	1.106	1.607	2.044	1.795	1.761	1.703	1.841	14.571	21.817	22.723	23.600
Gal9As19H42	0.109	0.0	3.158	101.233	11.712	11.782	10.504	10.647	9.527	0.0	0.0	0.0	0.0
S187H76	0.208	0.0	0.0	0.0	16.080	15.123	13.743	13.481	12.879	0.0	0.0	0.0	0.0
SiH4	0.005	0.046	0.090	0.544	0.801	0.978	1.059	0.664	1.002	5.807	6.400	8.067	10.537
apache1	0.035	0.094	0.750	1.231	1.549	1.407	1.416	1.482	1.535	12.918	21.983	30.059	22.131
apache2	0.142	0.840	5.206	11.462	9.348	7.725	6.123	6.535	6.231	113.364	0.0	0.0	0.0
auto	0.243	1.564	5.927	13.647	11.815	10.773	9.497	7.134	8.113	212.413	273.232	0.0	0.0
boyd2	0.074	0.134	1.364	62.858	4.383	3.674	3.291	10.078	6.315	18.90	21.477	21.597	36.917
c-55	0.017	0.066	0.295	1.529	1.562	1.592	1.849	1.399	1.720	17.803	36.937	48.612	80.473
c-70	0.031	0.094	0.766	1.627	2.068	2.268	2.307	1.443	1.984	18.863	26.485	36.458	48.367
ca-AstroPh	0.013	0.087	0.257	1.320	1.396	1.587	2.018	1.178	1.331	13.562	19.613	21.744	28.675
circuit_3	0.006	0.064	0.087	0.544	0.540	0.540	0.627	0.454	0.718	3.602	3.406	3.829	3.829
citationCiteSeer	0.099	1.542	3.264	29.040	9.516	8.943	7.717	7.221	5.452	252.821	0.0	0.0	0.0
copter2	0.023	0.151	0.733	1.269	1.924	1.819	1.861	1.547	1.734	14.522	30.263	30.062	27.043
crystk03	0.025	0.084	0.210	0.690	2.424	2.507	2.788	1.644	2.396	16.430	15.065	15.106	17.250
delaunay_n20	0.257	1.370	6.845	11.794	12.832	10.401	9.045	7.237	7.808	127.681	214.854	270.887	252.652
dictionary28	0.013	0.059	0.343	0.729	1.091	1.164	1.269	1.258	1.330	9.888	7.965	7.783	8.655
ex10	0.006	0.012	0.030	0.048	0.533	0.566	0.584	0.254	0.639	3.520	3.189	3.583	3.347
gupta2	0.065	2.075	0.900	0.0	7.038	8.940	7.081	6.420	5.357	0.0	0.0	0.0	0.0
gupta3	0.116	0.688	0.786	0.0	10.923	11.779	11.283	7.724	7.261	0.0	0.0	0.0	0.0
lpl1	0.011	0.080	0.373	1.668	1.149	1.206	1.481	1.377	1.413	6.590	8.509	8.016	11.014
mempplus	0.006	0.032	0.073	0.183	0.674	0.657	0.832	0.424	0.936	4.374	5.776	6.287	8.373
nasasrb	0.064	0.169	0.613	1.539	3.587	3.526	2.807	2.406	3.199	20.019	20.685	20.413	21.757
ncvxbqp1	0.017	0.068	0.420	0.604	1.238	1.154	1.169	1.134	1.206	8.660	17.101	11.923	11.470
parabolic_fem	0.088	0.348	2.937	5.803	6.332	5.607	5.099	4.100	4.370	44.500	63.501	61.983	95.454
rajat6	0.004	0.009	0.061	0.078	0.532	0.555	0.607	0.408	0.779	2.802	3.827	4.136	4.066
rajat10	0.018	0.019	0.148	0.204	0.730	0.727	0.767	0.631	0.925	4.317	4.326	4.974	5.796
rajat15	0.019	0.078	0.399	1.008	1.330	1.268	1.355	1.495	1.486	9.842	10.911	14.023	17.835
airpoll_conditional_st3	0.025	0.085	0.318	12.282	2.324	2.515	2.203	2.252	2.203	53.487	43.843	45.646	49.021
airpoll_prior_st3	0.017	0.084	0.150	1.830	2.327	2.544	2.229	1.309	2.229	37.342	35.457	35.548	36.613
cnt_cp2k	0.064	0.277	0.717	—	6.220	7.732	5.158	4.865	5.158	0.0	0.0	0.0	0.0
cnt_w90	0.005	0.010	0.148	0.204	0.730	0.727	0.767	0.631	0.925	4.317	4.326	4.974	5.796
kmc_potential_0	0.332	1.094	5.709	61.167	13.940	13.016	10.575	11.459	10.575	0.0	0.0	0.0	0.0
kmc_potential_1	1.515	0.0	0.0	317.549	55.548	49.505	34.982	37.262	34.982	0.0	0.0	0.0	0.0
qubit_fem_b4	0.068	0.453	2.193	22.243	4.611	4.305	4.278	4.672	4.278	237.962	72.351	0.0	0.0
sinw_w90	0.097	0.235	0.269	5.145	4.872	5.167	3.668	2.223	3.668	0.0	0.0	0.0	0.0

Table A.2.: Reordering time (seconds) for each matrix and reordering algorithm.

## A. Experimental Results

Matrix	RCM	AMD	METIS	NDCustom	ParMETIS-2	ParMETIS-4	ParMETIS-8	SCOTCH	PT-SCOTCH	db-2	db-4	db-8	db-16
144	3357	3357	55668	149388	151504	151220	150680	82500	85588	12	12	12	12
3Dspectralwave	1150k	0	1121k	0	2750k	2750k	2014k	2011k	0	0	0	0	0
598a	900	536	41088	127660	130220	129460	128680	80972	81096	292	412	344	616
CO	11978	0	11978	608508	580512	520376	609568	342212	342204	0	0	0	0
G2_circuit	1135	1135	4864	13312	16072	15480	14840	9776	10332	48	12	12	64
Gal9As19H42	151000	0	12520	643008	679272	678784	643800	445332	445008	0	0	0	0
S187H76	159656	0	0	0	824620	822872	781728	579816	5811828	0	0	0	0
SiH4	269	269	847	847	4	4	4	4	4	4	12	12	12
apache1	847	847	7528	7528	373028	375932	337520	374576	246656	249608	63816	0	0
apache2	7528	7528	10358	10358	52376	576724	525164	524616	315020	313500	53720	12	0
auto	10358	10358	44116	44116	9824	101928	100894	100392	102548	103336	108	12	104
boyd2	2344	2344	630	630	4312	4556	4668	4728	4	4	64	12	12
c-55	630	630	1030	1030	4	4	4	4	4	4	12	12	12
c-70	1030	1030	619	619	13620	13672	13884	13404	4128	3384	12	12	12
ca-AstroPh	619	619	69	69	4	4	4	4	4	4	12	12	12
circuit3	69	69	3615	3615	35420	154172	156948	155652	155100	101168	101444	12	0
citationCiteSeer	1187	1187	1187	1187	4	4	4	4	4	4	12	0	0
copter2	2736	2736	10080	10080	101108	101100	114396	52984	52988	1652	1064	1104	2764
crystk03	9830	9830	1528	637464	665980	640752	639516	490916	442864	75320	136	81396	296
delanunay_n20	278	278	5728	5728	5728	7824	7040	4668	4968	16	16	16	12
dictionary28	86	86	9992	6638	6638	86	4	4	4	4	12	12	12
ex10	173344	145608	139524	0	280616	247160	313728	171140	137372	0	0	0	0
guptia2	547	547	547	547	247344	249628	249176	277156	164000	169548	417288	0	0
guptia3	5742	5742	19036	247344	73	73	4	4	4	4	12	12	12
lpl1	513	513	513	513	4	4	4	4	4	4	12	12	12
memplus	158	158	4183	4183	124060	123856	123856	123776	63288	63748	12	12	12
nasastib	693	693	693	693	4	4	4	4	4	4	12	12	12
ncvxbqp1	1904	1904	1904	1904	69580	69920	69924	31688	31820	31688	368	268	816
parabolic_fem	1884	1884	1884	1884	63548	63796	63740	25948	26368	25948	12	460	280
rajat06	30944	9616	9616	0	383284	432120	236364	188276	236364	0	0	0	0
rajat10	204	204	204	204	4	4	4	4	4	4	12	12	12
rajat15	693	693	693	693	4	4	4	4	4	4	12	12	12
airpoll_conditional_st3	144440	78184	109184	899528	901560	938900	627736	668216	668216	0	0	0	0
airpoll_prior_st3	921224	0	0	3863608	3863608	3862000	2778032	2773948	2778032	0	0	0	0
cnt_cp2k	43616	43616	24760	193464	193972	233052	109816	109204	109816	35500	76532	0	0
qubit_fem_b4	76128	46336	48448	417968	417920	430000	249640	249648	249648	0	0	0	0
sinvw_w90	112	112	2724	3240	2984	248	248	368	368	1768	1128	1588	1588

Table A.3.: Maximum memory usage (KB) during reordering for each matrix and algorithm.

## A. Experimental Results

Matrix	RCM	AMD	METIS	NDCustom	ParMETIS-2	ParMETIS-4	ParMETIS-8	SCOTCH	PT-SCOTCH	db-2	db-4	db-8	db-16
144	2350	27834	24469	40703	58832	48525	12380	16450	12.0	12.0	12.0	12.0	12.0
3Dspectralwave	461793	0	460086	0	1120933	1317370	425597	548462	0.0	0.0	0.0	0.0	0.0
598a	472	211	13696	22727	46393	50232	49445	16556	16260	292	411	305	615
CO	8385	0	8385	48144	175532	155282	219679	50329	58432	0.0	0.0	0.0	0.0
G2_circuit	795	2432	1784	3398	3839	3803	335	335	48	12	12	12	64
Ga19As19t42	58637	0	60509	19684	260046	265948	245076	87241	98637	0.0	0.0	0.0	0.0
S187H76	60738	0	0	0	280158	300024	298703	121061	123001	0.0	0.0	0.0	0.0
SiH4	188	188	188	3	3	3	3	3	3	12	12	12	12
apache1	593	593	593	3	3	3	3	3	3	12	12	12	12
apache2	5270	5270	5270	26732	91929	94912	125461	34939	37490	70	0.0	0.0	0.0
auto	7251	7251	54253	198083	189256	206590	67046	71388	31	12	0.0	0.0	0.0
boyd2	1641	1641	2376	3267	24068	19810	23174	5428	6840	108	12	104	12
c-55	441	441	441	254	680	631	537	3	3	64	12	12	12
c-70	721	721	721	3	3	3	3	3	3	12	12	12	12
ca-AstroPh	433	433	433	1865	2634	3039	2391	1144	401	12	12	12	12
circuit_3	48	48	48	3	3	3	3	3	3	12	12	12	12
citationCiteSeer	2530	2530	11807	9830	34397	34500	29824	7963	10522	12	0.0	0.0	0.0
copter2	831	831	831	3	3	3	3	3	3	12	12	12	12
crystk03	1915	1915	1915	50792	42837	42193	51024	18628	11867	1520	15	1100	294
delanunay_n20	6881	6881	382	59146	199813	207967	242155	130707	102073	444	136	202	296
dictionary28	195	195	195	1460	2297	2268	2000	754	984	16	16	16	12
ex10	60	60	60	60	3	3	3	3	3	12	12	12	12
guptia2	4996	4647	4647	0	68475	98880	102083	24630	22404	0.0	0.0	0.0	0.0
guptia3	58665	86992	61466	0	272550	260201	231491	115501	109466	0.0	0.0	0.0	0.0
hpII	359	359	359	4	3	3	3	3	3	12	12	12	12
memplus	110	110	110	3	3	3	3	3	3	12	12	12	12
nasastb	2928	2928	2928	38015	50166	51248	38556	16725	12230	12	12	12	12
ncvxbqp1	383	383	383	3	3	3	3	3	3	12	12	12	12
parabolic_fem	4019	4019	9518	31722	64371	72374	108678	27836	28283	587	300	34	304
raja06	51	51	51	2	3	3	3	3	4	12	12	12	12
rajat10	143	143	143	3	3	3	3	3	3	12	12	12	12
rajat15	485	485	485	3	3	3	3	3	3	12	12	12	12
airpoll_conditional_st3	1333	1333	1333	6057	24380	22278	6247	6186	6247	368	252	643	799
airpoll_prior_st3	1319	1319	1319	28778	22287	20688	4871	9500	4871	0	443	264	28
cnt_cp2k	15472	6731	6731	0	119047	145990	47787	38493	47787	0	0	0	0
cnt_w90	78	78	78	810	441	487	231	1	231	361	195	44	167
kmc_potential_0	43898	32282	63399	73563	321702	392351	160563	149829	160563	0	0	0	0
kmc_potential_1	438089	0	0	279053	1605689	1848255	869721	790627	869721	0	0	0	0
qubit_fem_b4	21808	2	12380	14494	66102	90042	19132	17100	19132	375	121	0	0
sinvw_w90	38064	15447	16151	144526	174394	171231	65498	122032	65498	0	0	0	0

Table A.4: Average memory usage (KB) during reordering for each matrix and algorithm.

## A. Experimental Results

Matrix	RCM	AMD	METIS	NDCustom	ParMETIS-2	ParMETIS-4	ParMETIS-8	SCOTCH	PT-SCOTCH	db-2	db-4	db-8	db-16
144	380	56	47	86	33	30	31	265	215	49	79	66	85
3Dspectralwave	569	830	279	830	244	254	258	966	—	830	830	830	830
598a	398	46	57	62	30	26	27	182	188	52	52	64	113
CO	982	—	160	839	—	—	—	1103	1070	—	—	—	—
G2_circuit	212	11	12	101	100	99	103	93	54	11	33	100	100
Ga19As19H42	—	—	639	560	—	—	—	—	1009	—	—	—	—
S187H76	1012	—	—	—	—	—	—	1039	929	—	—	—	—
SiH4	329	382	130	397	220	192	165	194	399	173	468	609	798
apache1	206	4	12	12	106	120	110	22	35	19	25	13	18
apache2	357	2	11	13	305	318	332	131	85	17	603	603	603
auto	622	58	89	171	41	39	37	351	329	82	102	204	204
boyd2	5	4	3	2	6	10	9	3	5	3	3	3	3
c-55	10	23	37	26	31	30	26	33	25	29	27	23	19
c-70	10	19	42	36	22	32	25	23	27	33	29	29	26
ca-AstroPh	184	105	112	124	52	56	49	183	246	113	116	115	105
circuit_3	19	27	12	14	15	12	13	17	17	12	13	12	11
citationCiteseer	76	40	47	52	30	36	30	78	133	48	32	32	32
copt2	404	60	53	57	41	42	43	148	91	65	76	79	82
crypt03	719	251	126	125	212	273	164	290	328	173	179	203	218
delanunay_n20	—	24	26	36	42	37	37	—	—	29	25	24	25
dictionary28	67	35	30	35	24	30	33	70	46	37	29	30	36
ex10	354	92	55	67	53	47	61	49	59	73	59	79	62
gupta2	83	39	59	26	87	81	99	55	54	26	26	26	26
gupta3	531	329	225	301	953	788	470	572	668	301	301	301	301
hp11	33	20	29	21	21	22	23	23	21	29	20	25	21
memplus	34	33	48	40	25	14	19	45	46	47	43	34	35
nasasib	—	447	166	135	152	125	169	327	487	171	236	280	324
ncvxbqp1	118	18	22	22	16	14	15	95	943	23	23	26	25
parabolic_fem	—	31	22	24	17	16	17	964	743	23	23	26	25
raja06	81	5	7	6	31	29	24	10	10	7	6	6	7
raja10	135	5	6	6	27	35	33	10	10	7	6	8	8
raja15	77	81	62	66	31	33	35	92	99	60	56	59	58
airpoll_conditional_st3	360	400	330	422	243	291	817	1096	817	659	690	473	514
airpoll_prior_st3	428	452	390	358	223	235	775	806	775	676	611	517	760
cnt_cp2k	—	—	—	1609	—	—	—	—	—	—	—	—	—
cnt_w90	767	735	95	111	190	223	767	223	95	95	95	95	95
kmc_potential_0	—	232	114	237	202	201	427	451	427	555	555	555	555
kmc_potential_1	—	571	287	383	332	810	706	810	571	571	571	571	571
qubit_fem_b4	248	116	45	60	82	75	140	104	140	98	80	225	225
sinw_w90	—	—	910	908	—	—	910	961	910	—	—	—	—

Table A.5.: Elimination tree depth for each matrix and reordering algorithm.

# Appendix **B**

## Task Description

Include the task description PDF file you got from your advisors with  
`\includepdf[pages=-, scale=0.9]{../path/to/task/description.pdf}`.

# List of Acronyms

ASIC	application-specific integrated circuit
DFT	design for testability
DZ	Microelectronics Design Center at ETH Zurich
FPGA	field-programmable gate array
IC	integrated circuit
IIS	Integrated Systems Laboratory
PDF	Portable Document Format
SNR	signal-to-noise ratio
SPSE	Situation-Problem-Solution-Evaluation
SVG	scalable vector graphics
VCS	version control system
WYSIWYG	“what you see is what you get”

# List of Figures

1.1.	The bcsstk32 matrix from automotive chassis analysis. This $28,924 \times 28,924$ symmetric matrix represents an automobile chassis structure analyzed using finite element modeling. . . . .	1
1.2.	Sparse matrix structure from a quantum transport device. The block-tridiagonal pattern reflects local physical interactions between neighboring spatial layers. . . . .	2
2.1.	Illustration of elimination tree construction for the given sparse symmetric matrix. The sparsity is denoted by filled black circles and the fill-in induced is denoted by hollow red circles. . . . .	7
2.2.	Separator identification process. The separator vertices divide the graph into approximately equal subgraphs . . . . .	11
4.1.	Fill-in ratio (fill-in/original nnz) comparison across different matrix categories . . . . .	27
4.2.	Reordering time scaling with matrix size (number of non-zeros) for different algorithms. Each subplot shows the relationship between matrix size and computational time for various reordering methods. . . . .	29
4.3.	Sparsity patterns of matrix 144 with different reordering algorithms applied. The natural ordering (top-left) shows a scattered pattern leading to high fill-in, while optimized reorderings (other panels) demonstrate more defined structures. . . . .	31
4.4.	Similarly, sparsity patterns of matrix copter2 with different reordering algorithms applied. . . . .	32
4.5.	Elimination tree depth comparison across different matrix categories for various reordering algorithms. Lower depth indicates better parallelization potential. Blue indicates lower depth (better), while pink indicates higher depth (worse). . . . .	33
4.6.	Speedup of GPU-accelerated RCM compared to CPU-based RCM on large matrices (over 10K nodes). . . . .	34

*List of Figures*

5.1. Recursive calls in nested dissection . . . . .	36
---	----

# List of Tables

4.1.	Fritz cluster node configuration . . . . .	22
4.2.	SuiteSparse matrix dataset used for evaluation . . . . .	24
4.3.	Quantum Transport matrix dataset used for evaluation . . . . .	24
5.1.	GNN-based method performance comparison with RCM and MinDegree	36
A.1.	Fill-in (number of nonzeros after factorization) for each matrix and reordering algorithm. . . . .	40
A.2.	Reordering time (seconds) for each matrix and reordering algorithm. . . . .	41
A.3.	Maximum memory usage (KB) during reordering for each matrix and algorithm. . . . .	42
A.4.	Average memory usage (KB) during reordering for each matrix and algorithm. . . . .	43
A.5.	Elimination tree depth for each matrix and reordering algorithm. . . . .	44

# Bibliography

- [1] M. Yannakakis, "Computing the Minimum Fill-In is NP-Complete," *SIAM Journal on Algebraic Discrete Methods*, vol. 2, no. 1, pp. 77–79, Mar. 1981, publisher: Society for Industrial and Applied Mathematics. [Online]. Available: <https://epubs.siam.org/doi/10.1137/0602010>
- [2] R. J. Lipton, D. J. Rose, and R. E. Tarjan, "Generalized Nested Dissection," *SIAM Journal on Numerical Analysis*, vol. 16, no. 2, pp. 346–358, Apr. 1979, publisher: Society for Industrial and Applied Mathematics. [Online]. Available: <https://epubs.siam.org/doi/10.1137/0716027>
- [3] A. George, "Nested Dissection of a Regular Finite Element Mesh," *SIAM Journal on Numerical Analysis*, vol. 10, no. 2, pp. 345–363, Apr. 1973, publisher: Society for Industrial and Applied Mathematics. [Online]. Available: <https://epubs.siam.org/doi/10.1137/0710032>
- [4] U. V. Çatalyürek, C. Aykanat, and E. Kayaaslan, "Hypergraph Partitioning-Based Fill-Reducing Ordering for Symmetric Matrices," *SIAM Journal on Scientific Computing*, vol. 33, no. 4, pp. 1996–2023, Jan. 2011, publisher: Society for Industrial and Applied Mathematics. [Online]. Available: <https://epubs.siam.org/doi/10.1137/090757575>
- [5] A. Dasgupta and P. Kumar, "Alpha Elimination: Using Deep Reinforcement Learning to Reduce Fill-In During Sparse Matrix Decomposition," in *Machine Learning and Knowledge Discovery in Databases: Research Track: European Conference, ECML PKDD 2023, Turin, Italy, September 18–22, 2023, Proceedings, Part IV*. Berlin, Heidelberg: Springer-Verlag, Sep. 2023, pp. 472–488. [Online]. Available: [https://doi.org/10.1007/978-3-031-43421-1\\_28](https://doi.org/10.1007/978-3-031-43421-1_28)
- [6] Y.-H. Chang, A. Buluç, and J. Demmel, "Parallelizing the approximate minimum degree ordering algorithm: Strategies and evaluation," 2025. [Online]. Available: <https://arxiv.org/abs/2504.17097>
- [7] D. Merrill, "Scalable GPU Graph Traversal."
- [8] P. Kaleta, "kaletap/bfs-cuda-gpu," May 2025, original-date: 2019-11-14T12:55:50Z. [Online]. Available: <https://github.com/kaletap/bfs-cuda-gpu>

## Bibliography

- [9] C. Yuan, "Fast Sparse Matrix Reordering on GPU for Cholesky Based Solvers."
- [10] M. S. Gilbert, K. Madduri, E. G. Boman, and S. Rajamanickam, "Jet: Multilevel Graph Partitioning on Graphics Processing Units," *SIAM Journal on Scientific Computing*, vol. 46, no. 5, pp. B700–B724, Oct. 2024, publisher: Society for Industrial and Applied Mathematics. [Online]. Available: <https://pubs.siam.org/doi/10.1137/23M1559129>
- [11] H. Kaeslin, *Digital Integrated Circuit Design: From VLSI Architectures to CMOS Fabrication*. Cambridge University Press, Apr. 2008.
- [12] D. Hankerson, S. Vanstone, and A. Menezes, *Guide to Elliptic Curve Cryptography*, ser. Springer Professional Computing. Springer, 2004.
- [13] NIST, *Advanced Encryption Standard (AES) (FIPS PUB 197)*, National Institute of Standards and Technology, Nov. 2001.
- [14] P. Rogaway, M. Bellare, J. Black, and T. Krovetz, "OCB: A block-cipher mode of operation for efficient authenticated encryption," in *ACM Conference on Computer and Communications Security*, 2001, pp. 196–205.
- [15] Xilinx. (2011, Nov.) Virtex-6 FPGA Configuration User Guide. UG360 (v3.4). [Online]. Available: [http://www.xilinx.com/support/documentation/user\\_guides/ug360.pdf](http://www.xilinx.com/support/documentation/user_guides/ug360.pdf)
- [16] ——. (2011, Oct.) 7 Series FPGAs Configuration User Guide. UG470 (v1.2). [Online]. Available: [http://www.xilinx.com/support/documentation/user\\_guides/ug470\\_7Series\\_Config.pdf](http://www.xilinx.com/support/documentation/user_guides/ug470_7Series_Config.pdf)
- [17] Wikipedia, "Isaac Newton," accessed October 1, 2012. [Online]. Available: [http://en.wikipedia.org/w/index.php?title=Isaac\\_Newton&oldid=514997436](http://en.wikipedia.org/w/index.php?title=Isaac_Newton&oldid=514997436)
- [18] D. E. Knuth, *The TeXbook*. Addison-Wesley, 1984.
- [19] S. Pakin, *The Comprehensive LaTeX Symbol List*, Jan. 2017. [Online]. Available: <http://mirrors.ctan.org/info/symbols/comprehensive/symbols-a4.pdf>
- [20] U. Kern, *xcolor: Driver-independent color extensions for L<sup>A</sup>T<sub>E</sub>X and pdfL<sup>A</sup>T<sub>E</sub>X*, May 2016. [Online]. Available: <http://mirrors.ctan.org/macros/latex/contrib/xcolor/xcolor.pdf>
- [21] S. M. Moser, *How to Typeset Equations in L<sup>A</sup>T<sub>E</sub>X*, Sep. 2017. [Online]. Available: [http://moser-isi.ethz.ch/docs/typeset\\_equations.pdf](http://moser-isi.ethz.ch/docs/typeset_equations.pdf)
- [22] J. Wright, *siunitx: A comprehensive (SI) units package*, Aug. 2017. [Online]. Available: <http://mirrors.ctan.org/macros/latex/contrib/siunitx/siunitx.pdf>
- [23] T. Oetiker, *The Not So Short Introduction to L<sup>A</sup>T<sub>E</sub>X2e*, Jun. 2016. [Online]. Available: <https://tobi.oetiker.ch/lshort/lshort.pdf>

## Bibliography

- [24] F. Mittelbach, M. Goossens, J. Braams, D. Carlisle, and C. Rowley, *The L<sup>A</sup>T<sub>E</sub>X Companion*, ser. Tools and Techniques for Computer Typesetting. Pearson Education, 2004.
- [25] F. Mittelbach, "How to influence the position of float environments like figure and table in L<sup>A</sup>T<sub>E</sub>X," *TUGboat*, vol. 35, no. 3, pp. 248–254, 2014. [Online]. Available: <https://www.latex-project.org/publications/tb111mitt-float.pdf>
- [26] G. Grätzer, *More Math Into L<sup>A</sup>T<sub>E</sub>X*. Springer International Publishing, 2016.
- [27] W. Strunk, *Elements of Style*. Start Publishing LLC, 2012.
- [28] U. of Chicago Press and U. of Chicago Press Editorial Staff, *The Chicago Manual of Style*. University of Chicago Press, 2017.
- [29] T. Miller and D. Parker, "Writing for the reader: A problem-solution approach," *English Teacher Forum*, no. 3, pp. 21–27, 2012. [Online]. Available: <http://files.eric.ed.gov/fulltext/EJ997525.pdf>
- [30] M. Hoey, *On the Surface of Discourse*. Allen and Unwin, 1983.
- [31] K. L. Turabian, *A Manual for Writers of Research Papers, Theses, and Dissertations*. University of Chicago Press, 2013.